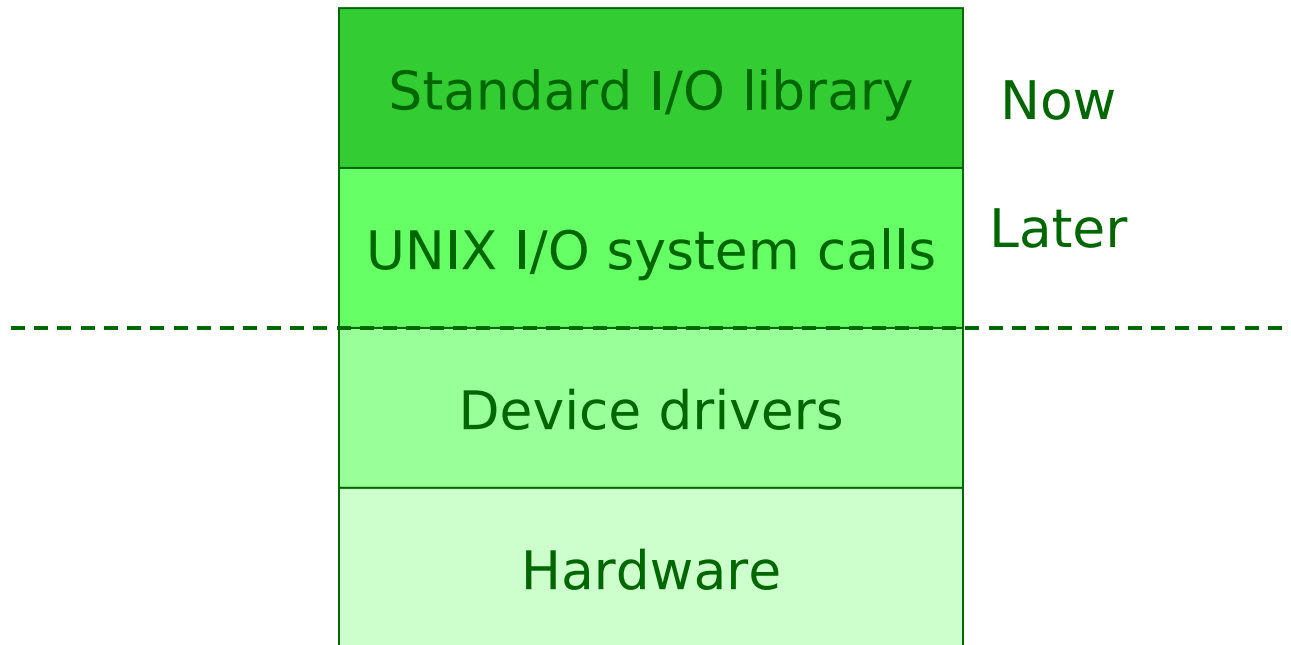


C's Standard I/O Library

Alan L. Cox
alc@cs.rice.edu

Layers of I/O Abstraction



Standard I/O vs. UNIX I/O:

- ♦ Only the common features
- ♦ Easier to use

Character Output

```
#include <stdio.h>
```

```
int
```

```
main(void)
```

```
{
```

```
    putchar('H');
```

```
    putchar('e');
```

```
    putchar('l');
```

```
    putchar('l');
```

```
    putchar('o');
```

```
    putchar('\n');
```

```
    return (0);
```

```
}
```

← Won't continually remind you of this

```
UNIX% ./hello
```

```
Hello
```

```
UNIX%
```

Cumbersome to output one character at a time

- ♦ Still useful in some situations

Difficult to format output

String Output

```
#include <stdio.h>

int
main(void)
{
    puts("Hello\n");

    return (0);
}
```

```
UNIX% ./hello
Hello
UNIX%
```

Better than a character at a time

- ♦ Still difficult to format output

Formatted Output

```
void  
hello(char *name, int hour, int min)  
{
```

```
    printf("Hello, %s, it's %d:%d.\n", name, hour, min);  
}
```

```
int  
main(void)  
{
```

```
    hello("Alan", 2, 55);  
  
    return (0);  
}
```

Conversion specifications indicate the type of argument

man 3 printf for complete list

```
UNIX% ./hello  
Hello, Alan, it's 2:55.  
UNIX%
```

Character Input

```
#include <ctype.h>
#include <stdio.h>
/* Convert input to lowercase */
int
main(void)
{
    int c;

    c = getchar();
    while (c != EOF) {
        c = isupper(c) ? tolower(c) : c;
        putchar(c);
        c = getchar();
    }
    return (0);
}
```

Why int instead of char?

**EOF is not a character,
getchar returns this
value at the end-of-file**

**We could implement
these,
but they are part of the
standard library**

String Input

```
#include <stdio.h>
```

```
/* Echo */
```

```
int
```

```
main(void)
```

```
{
```

```
    char input[100];
```

```
    while (gets(input) != NULL) {
```

```
        puts(input);
```

```
    }
```

```
    return (0);
```

```
}
```

**What if input string > 99
characters? (+
terminator)
Don't use gets!**

**Returns NULL at the end-of-
file**

Formatted Input

```
int
main(void)
{
    char    name[50];
    int     num;
    float   rate;
    int     dollars, cents;

    scanf("%d %s %f", &num, name, &rate);

    dollars = (int)rate;
    cents   = (int)((rate - (float)dollars) * 100);

    printf("%s is employee #d and makes $d.%02d per hour.\n",
           name, num, dollars, cents);

    return (0);
}
```

What if input string > 49
characters? (specifications as
printf terminator)

Don't use scanf with
strings!

Pointers – Why?

Why not &name?

Ugly C
typecasting

More
formatting

Input: 36 Alan 90.5E-1

Output: Alan is employee #36 and makes \$9.05 per hour.

UNIX Shell Redirection

Can redirect keyboard & screen I/O to/from files

- ♦ I/O works with keyboard and screen by default
- ♦ Use redirection to read from/print to a file

```
UNIX%  program < infile
```

```
UNIX%  program > outfile
```

```
UNIX%  program >& outerrfile
```

```
UNIX%  program < inputfile > outputfile
```

```
other options depend on shell
```

What if you want to read/write multiple files?

Basic File Operations: Open & Close

```
FILE *fopen(const char *filename,  
            const char *mode);
```

FILE * = file “handle”

“r” – read
“w” – write
“a” – append
...

```
int fclose(FILE *stream);
```

0 iff successful

Open files automatically closed upon exit

Built-in Files

stdin – keyboard input

stdout – screen output – normal output

stderr – screen output – error output

No need to open/close

Why two?

Writing to Files

```
int putc(int c, FILE *stream);
```

- returns EOF if error
- putchar specialized to stdout

```
int fputs(const char *s, FILE *stream);
```

- returns the # of chars written (or EOF)
- puts specialized to stdout

```
int fprintf(FILE *stream,  
            const char *format,  
            ... );
```

- returns the # of chars written
- returns negative value on error
- printf specialized to stdout

Reading from Files

```
int getc(FILE *stream);
```

- returns character (or EOF)
- getchar specialized to stdin

```
char *fgets(char *s, int n, FILE *stream);
```

- returns s (or NULL)
- gets specialized to stdin
- max length argument (n) makes this safe!

```
int fscanf(FILE *stream,  
           const char *format,  
           ... );
```

- returns the # of matched and assigned inputs
- returns EOF on error
- scanf specialized to stdin

End of File

**Need a way to detect when everything
in a file has been read**

```
int feof(FILE *stream);
```

True/false

```
int c = getc(FILE *stream);  
if (c == EOF) ...
```

Not a char

Example: Character I/O

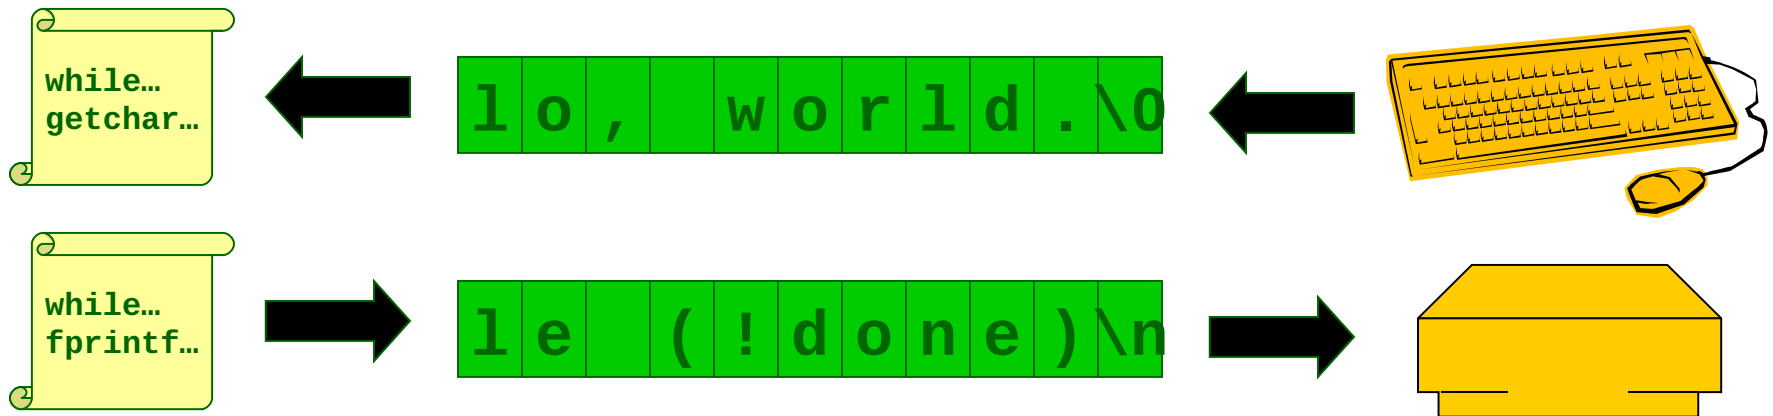
```
FILE *infile  = fopen(infilename, "r");
FILE *outfile = fopen(outfilename, "w");
int c;

/* Verify that fopen() succeeded. */
while (!feof(infile) {
    c = getc(infile);
    putc(c, outfile);
}
fclose(infile);
fclose(outfile);
```

Buffering

Streams are buffered

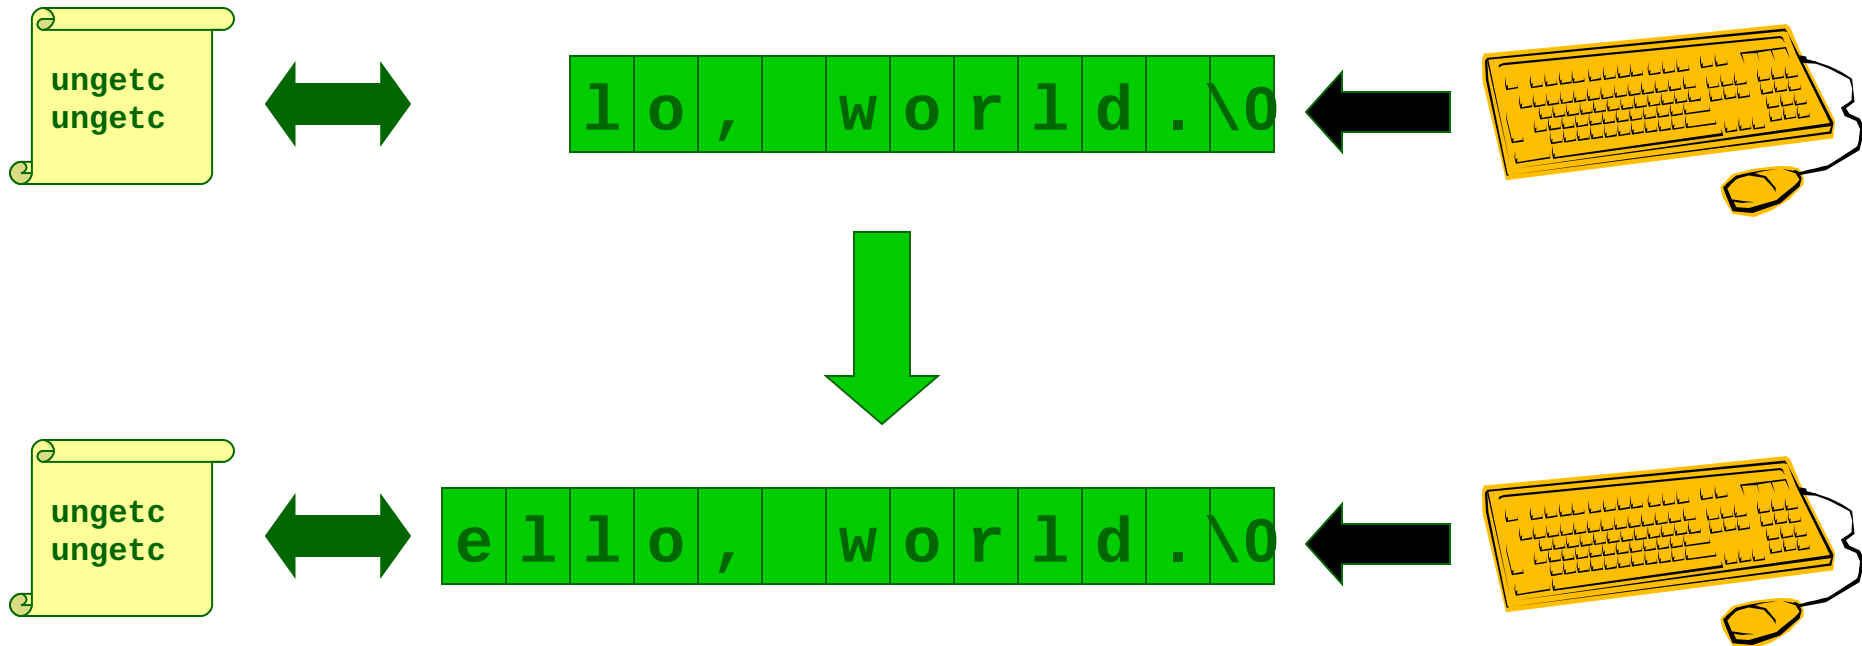
Manages data being produced & consumed at different rates
Transfer data in blocks, for efficiency



Main exception: **stderr** unbuffered

Buffering

```
int ungetc(int c, FILE *stream);
```



Buffering & Flushing

***Flush* = empty buffer to output file**

When:

- ♦ Internal buffer full
- ♦ Closing a file
- ♦ If line-buffered (e.g., stdout) → writing newline

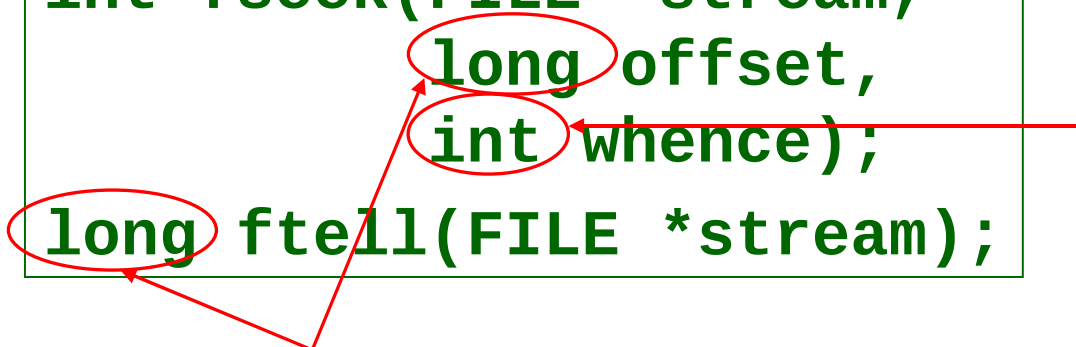
Mainly important in case of errors/interrupts

```
int fflush(FILE *stream);
```

Seeking

Can also write/read via random access:

```
int fseek(FILE *stream,  
          long offset,  
          int whence);  
long ftell(FILE *stream);
```



offset

Encodes “from here”, “from beginning”, or “from end”

Or if a long isn't big enough:

```
int fgetpos(FILE *stream, fpos_t *pos);  
int fsetpos(FILE *stream, const fpos_t *pos);
```

“Writing” & “Reading” To/From a String

```
int sprintf(char *s,  
            const char *format,  
            ...);  
int snprintf(char *s,  
            size_t n,  
            const char *format,  
            ...);  
int sscanf(const char *s,  
            const char *format,  
            ...);
```

What's Wrong With This Code?

```
char  *s = "Hello";  
  
gets(s);
```

Doesn't check maximum string size.

- E.g., if user enters "**123456789**", writes unallocated memory
- **Never** use **gets()** – use **fgets()** or repeated **getc()**

Basis for classic *buffer overflow attacks*:

- 1988 Internet worm
- Modern attacks on Web servers
- If overflows a stack-allocated buffer, can change return

What's Wrong With This Code?

```
int  value;  
char s[8];  
...  
scanf("%d %s", value, s);
```

&value



scanf() expects pointers

(Be aware when functions take pointer args!)

Like **gets()**, could overflow buffer

Don't use **scanf()** for strings

Writing & Reading Words

```
int putw(int w, FILE *stream);  
int getw(FILE *stream);
```

Compare:

```
putw(514, file1);  
fprintf(file2, "%d", 514);
```

? Results? ?

file1:

0x00	0x00	0x01	0x02
------	------	------	------

assumes big-endian

514 = 00000000000010010

file2:

0x35	0x31	0x33	0x00
------	------	------	------

'5' '1' '4' '\0'

Writing & Reading Bytes

```
size_t fwrite(const void *ptr,  
              size_t size,  
              size_t nitems,  
              FILE *stream);  
size_t fread(void *ptr,  
             size_t size,  
             size_t nitems,  
             FILE *stream);
```


More details

I have only outlined the functions' purposes
There are a handful of other functions
See reference book, man pages, or web for details

Next Time

Memory Allocation