

Assembly Language

Alan L. Cox
`alc@cs.rice.edu`

Why Learn Assembly Language?

You'll probably never write a program in assembly

- ♦ **Compilers are much better and more patient than you are**

But, understanding assembly is key to understanding the machine-level execution model

- ♦ **Behavior of programs in presence of bugs**
 - High-level language model breaks down
- ♦ **Tuning program performance**
 - Understanding sources of program inefficiency
- ♦ **Implementing system software**
 - Compiler has machine code as target
 - Operating systems must manage process state

Instruction Set Architecture

Contract between programmer and the hardware

- ♦ Defines visible state of the system
- ♦ Defines how state changes in response to instructions

Assembly Programmer (compiler)

- ♦ ISA is model of how a program will execute

Hardware Designer

- ♦ ISA is formal definition of the correct way to execute a program

Architecture vs. Implementation

Instruction Set Architecture

- ♦ Defines what a computer system does in response to a program and a set of data
- ♦ Programmer visible elements of computer system

Implementation

- ♦ Defines how a computer does it
- ♦ Sequence of steps to complete operations
- ♦ Time to execute each operation
- ♦ Hidden “bookkeeping” functions

Often Many Implementations of an ISA

ISA	Implementations
x86-64	Intel Core i7
	AMD Phenom II
	VIA Nano
SPARC V.9	UltraSPARC III
	HyperSPARC

Why separate architecture and implementation?

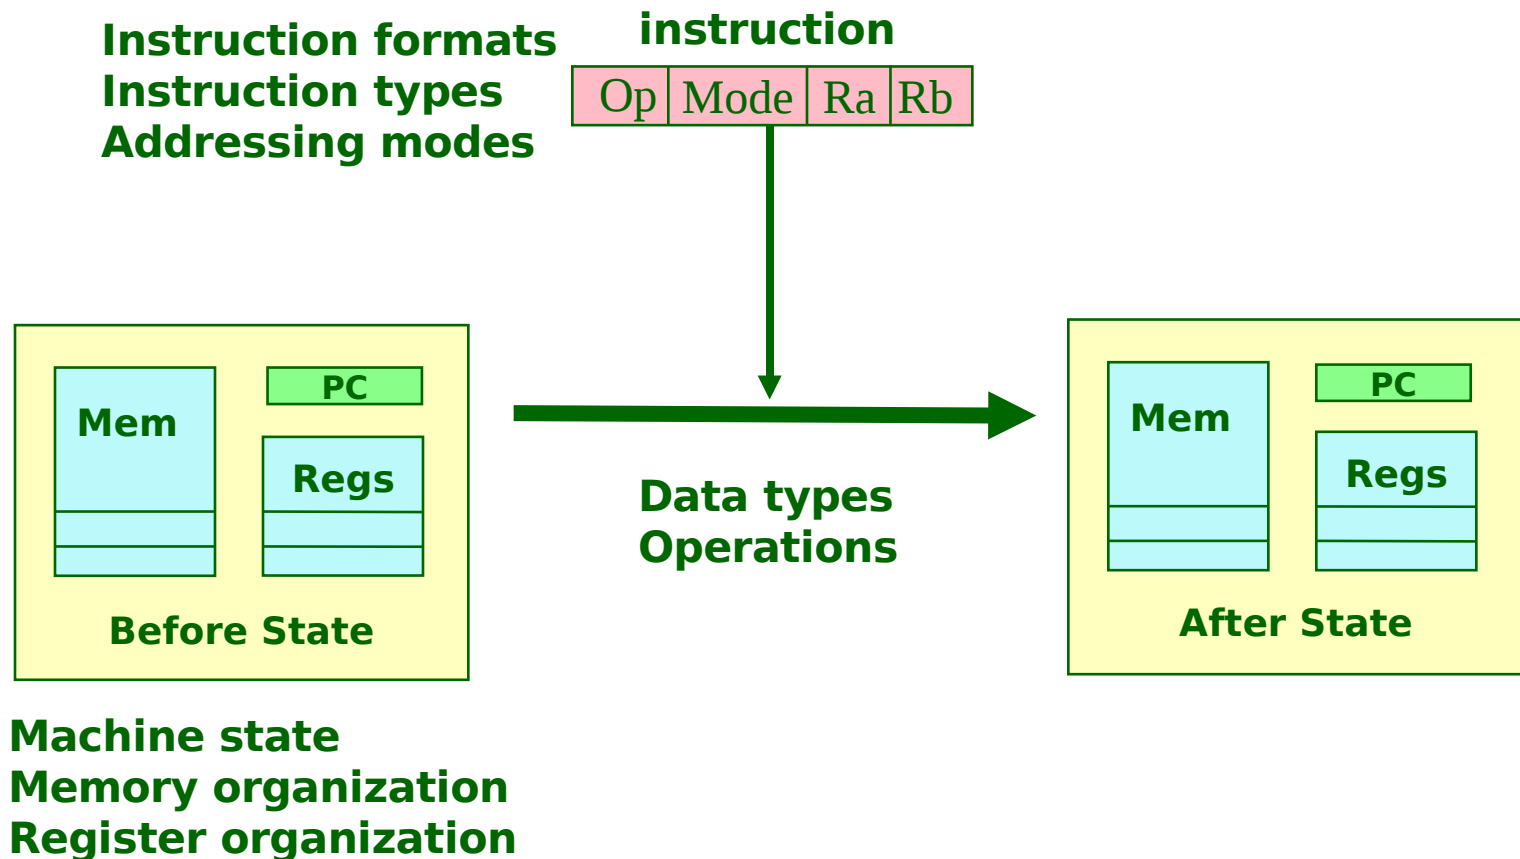
Compatibility

- ♦ **VAX architecture: mainframe \Rightarrow single chip**
- ♦ **ARM: 20x performance range**
 - high vs. low performance, power, price

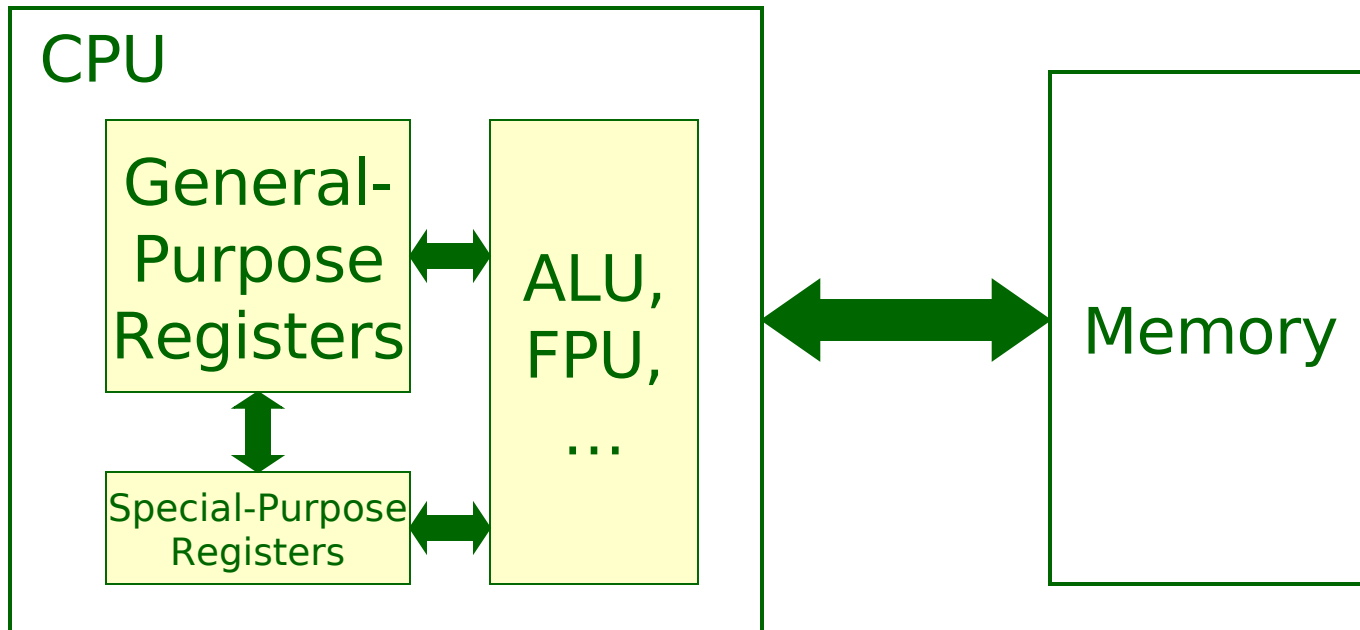
Longevity

- ♦ **20-25 years of ISA**
- ♦ **x86/x86-64 in 10th generation of implementations (architecture families)**
- ♦ **Retain software investment**
- ♦ **Amortize development costs over multiple markets**

Instruction Set Basics




Typical Machine State



x86-64 Machine State

General-purpose registers

- ♦ **16 64-bit registers**
 - **Extended x86:** %rax, %rbx, %rcx, %rdx,
 %rsp, %rbp, %rdi, %rsi
 - **New:** %r9 – %r15
- ♦ **%rsp is always the stack pointer**
- ♦ **Backward compatibility for extended x86 registers**



x86-64 Machine State

Special-purpose registers

- ♦ **16 128-bit floating point registers**
 - **SSE: %xmm0 - %xmm15**
- ♦ **Each register stores two double precision or four single precision values**
 - **Vector arithmetic for speed!**

x86-64 Machine State

Memory

- ♦ **64-bit addresses**
 - **However, only 48 bits used**
 - 47 bits of address space for programs
 - 47 bits of address for operating system
- ♦ **Byte addressed**
- ♦ **Little-endian byte ordering within a word**
 - **Least significant byte comes first in memory**

Assembly Language

One assembly instruction

- ♦ One group of machine language bits

But, assembly language has additional features:

- ♦ Distinguishes instructions & data
- ♦ Labels = names for program control points
- ♦ Pseudo-instructions = special directives to the assembler
- ♦ Macros = user-definable abbreviations for code & constants

Instructions

What do these instructions do?

- ♦ Same kinds of things as high-level languages!

Arithmetic & logic

- ♦ Core computation

Data transfer

- ♦ Copying data between memory locations and/or registers

Control transfer

- ♦ Changing which instruction is next

Machine Language Encodings

Variable-width:

- ♦ 1 to dozens of bytes
- ♦ Flexible ISA
- ♦ CISC
- ♦ Harder hardware implementation
- ♦ x86/x86-64, VAX, ...

Fixed-width:

- ♦ 1 word, typically
- ♦ Limited ISA
- ♦ RISC
- ♦ Easier hardware implementation
- ♦ SPARC, MIPS, PowerPC, ...

Look at representative examples



Machine Language Instruction Formats

Different instructions have different kinds of operands

Can't encode all fields of all instructions within one word

- ♦ Use different formats for different instructions
- ♦ Different architectures use different formats
- ♦ Different architectures group different sets of instructions

SPARC:

- ♦ Format 1: 1 30-bit immediate (call with immediate address)
- ♦ Format 2: 1 reg & 1 condition & 1 22-bit immediate (branches)
- ♦ Format 3: 3 regs, or 2 regs & 1 13-bit immediate (arithmetic, loads)

Machine Language: Arithmetic

add %o2,%o3,%o1			0x9202800B
A variant of Format 3	op	math	10
	rd	%o1	01001
	op3	add	000000
	rs1	%o2	01010
	i	immediate?	0
	-	N/A	00000000
	rs2	%o3	01011

32 registers \Rightarrow 5 bits

Not all bits important
Only used if i=1

Machine Language: Memory Ops

ld [%o0+4], %o2			0xD4022004
A variant of Format 3	op	memory	11
	rd	%o2	01010
	op3	ld	000000
	rs1	%o0	01000
	i	immediate?	1
	simm13	(13 bits)	00000000000100

Machine Language: Calls

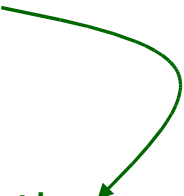
`call there`
`branch delay` ← Instruction after branch/jump is *always* executed!
...
...
`there: ...` ← Assume assembler/linker assigns **there** to address 0x20000000

call there			0x48000000
Format 1	op	call	01
	disp30	(30 bits)	00100000000000000000000000000000

call also saves the next program counter in %o7, which is used by ret/retl to return from a procedure (ret uses %i7, retl uses %o7)

Machine Language: Branches

branch to **dest**
branch delay
instruction
dest: destination instruction



3 instructions

be dest			0x02800003
Format 2	op	branch	00
	a	annul?	0
	cond	e	0001
	op2	b	010
	disp22	(22 bits)	00000000000000000000011

words relative
to the branch
instruction



Application Binary Interface (ABI)

Standardizes the use of memory and registers by C compilers

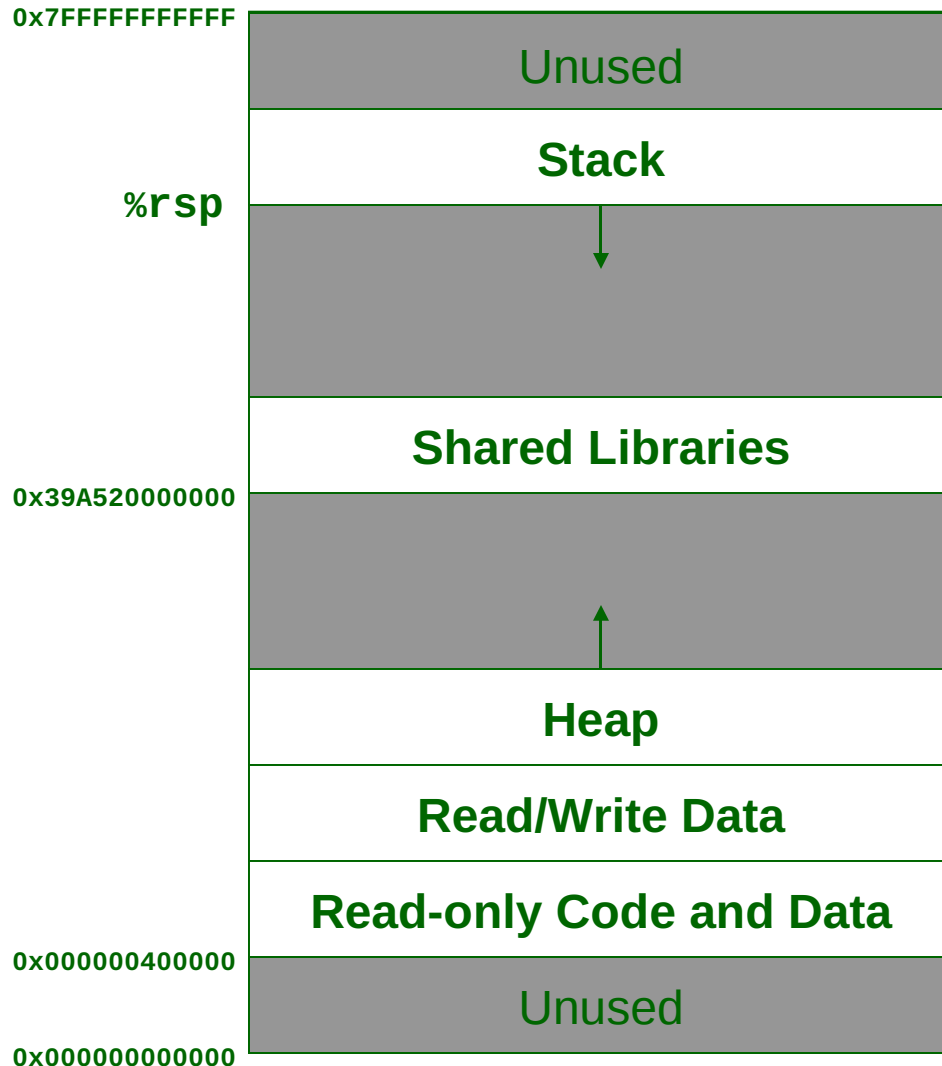
- ♦ **Enables interoperability of code compiled by different C compilers**
 - E.g., a program compiled with Intel's optimizing C compiler can call a library function that was compiled by the GNU C compiler
- ♦ **Sets the size of built-in data types**
 - E.g., int, long, etc.
- ♦ **Dictates the implementation of function calls**
 - E.g., how parameters and return values are passed

Register Usage

The x86-64 ABI specifies that registers are used as follows

- ♦ **Temporary (callee can change these)**
`%rax, %r10, %r11`
- ♦ **Parameters to function calls**
`%rdi, %rsi, %rdx, %rcx, %r8, %r9`
- ♦ **Callee saves (callee can only change these after saving their current value)**
`%rbx, %rbp, %r12-%r15`
 - `%rbp` is typically used as the “frame” pointer to the current function’s local variables
- ♦ **Return values**
`%rax, %rdx`

Procedure Calls and the Stack



Where are local variables stored?

- ♦ **Registers (only 16)**
- ♦ **Stack**

Stack provides as much local storage as necessary

- ♦ **Until memory exhausted**
- ♦ **Each procedure allocates its own space on the stack**

Referencing the stack

- ♦ **`%rsp` points to the bottom of the stack in x86-64**

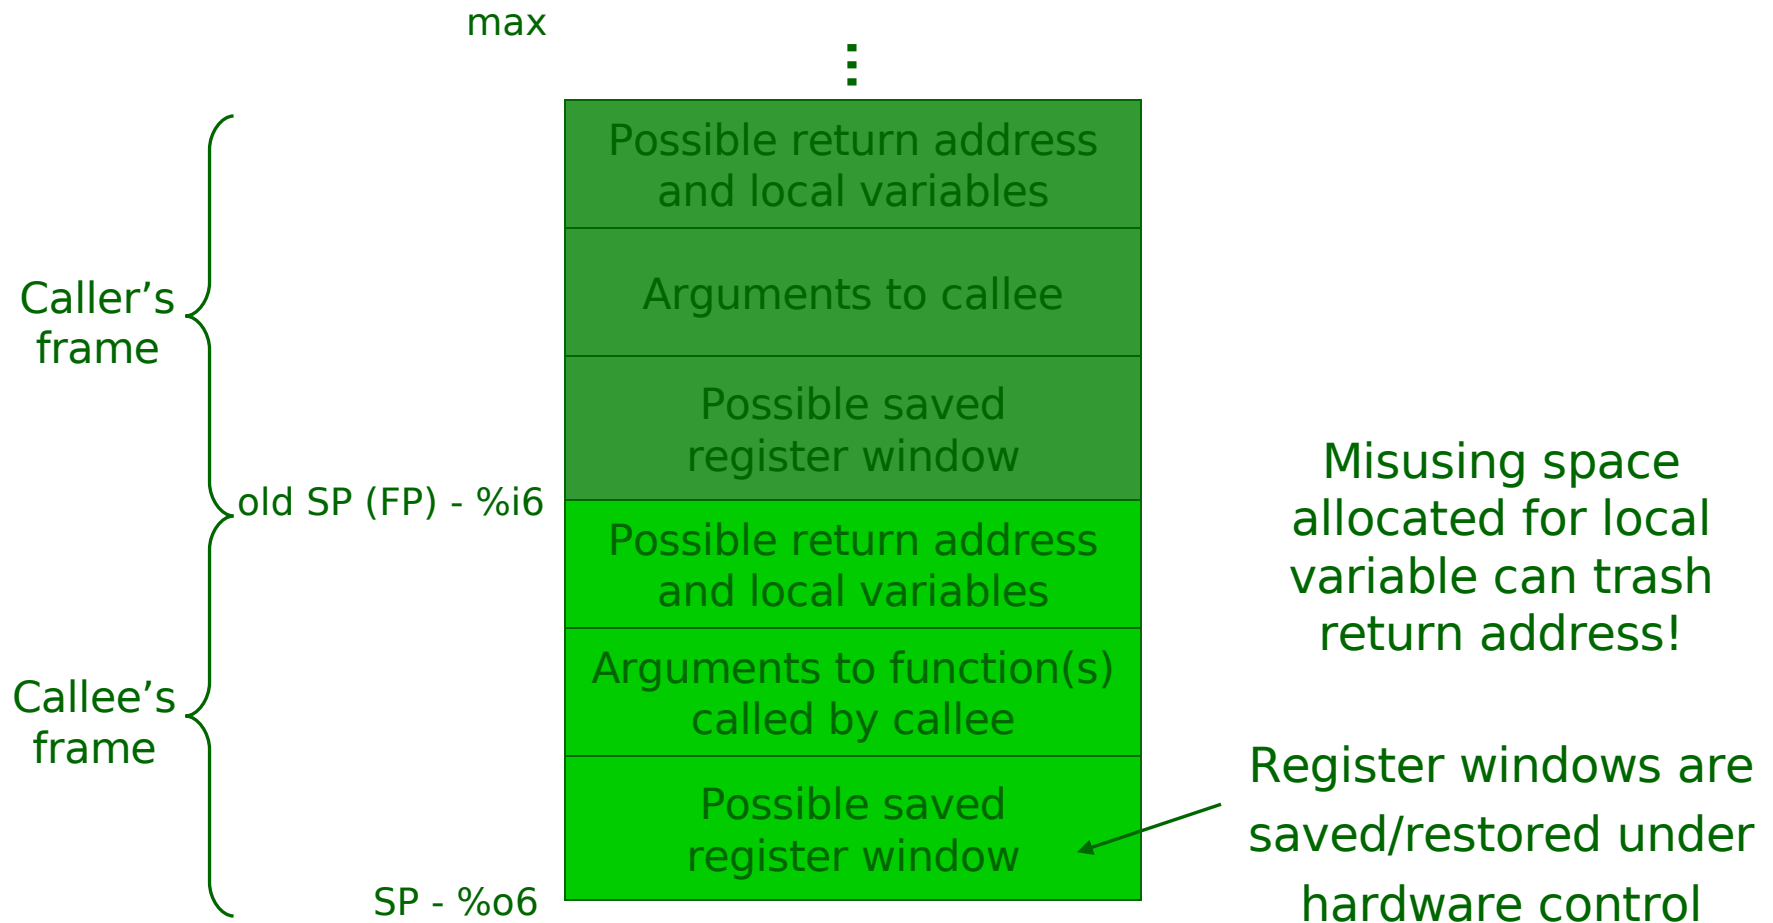
Control Flow: Function Calls

What must assembly/machine language do?

Caller	Callee
<ol style="list-style-type: none">1. Save function arguments2. Branch to function body	<ol style="list-style-type: none">3. Execute body<ul style="list-style-type: none">• May allocate memory• May call functions4. Save function result5. Branch to where called

2. Use `%eax` to store return value, `%ecx` to store error code

Program Stack



Example C Program

main.c:

```
#include <stdio.h>

void
hello(char *name, int hour, int min)
{
    printf("Hello, %s, it's %d:%02d.",
           name, hour, min);
}

int
main(void)
{
    hello("Alan", 2, 55);

    return (0);
}
```

Run the command:

UNIX% gcc -S main.c



Output a file named main.s
containing the assembly
code for main.c

C Compiler's Output

```
.file "main.c"
.section .rodata
.LC0:
.string "Hello, %s, it's %d:%02d."
.text
.globl hello
.type hello, @function
hello:
.LFB2:
pushq %rbp
.LCFI0:
movq %rsp, %rbp
.LCFI1:
subq $16, %rsp
.LCFI2:
movq %rdi, -8(%rbp)
movl %esi, -12(%rbp)
movl %edx, -16(%rbp)
movl -16(%rbp), %ecx
movl -12(%rbp), %edx
movq -8(%rbp), %rsi
movl $.LC0, %edi
movl $0, %eax
```

```
call printf
leave
ret
.LFE2:
.size hello, .-hello
.section .rodata
.LC1:
.string "Alan"
.text
.globl main
.type main, @function
main:
.LFB3:
pushq %rbp
.LCFI3:
movq %rsp, %rbp
.LCFI4:
movl $55, %edx
movl $2, %esi
movl $.LC1, %edi
call hello
movl $0, %eax
<..snip..>
```

Instructions: Opcodes

```
.file "main.c"
.section .rodata
.LC0:
.string "Hello, %s, it's %d:%02d."
.text
.globl hello
.type hello, @function
hello:
.LFB2:
    pushq    %rbp
.LCFI0:
    movq     %rsp, %rbp
.LCFI1:
    subq     $16, %rsp
.LCFI2:
    movq     %rdi, -8(%rbp)
    movl     %esi, -12(%rbp)
    movl     %edx, -16(%rbp)
    movl     -16(%rbp), %ecx
    movl     -12(%rbp), %edx
    movq     -8(%rbp), %rsi
    movl     $.LC0, %edi
    movl     $0, %eax
```

```
    call     printf
    leave
    ret
.LFE2:
    .size    hello, .-hello
<..snip..>
```

Arithmetic,
data transfer, &
control transfer

Instructions: Operands

```
.file "main.c"
.section .rodata
.LC0:
.string "Hello, %s, it's %d:%02d."
.text
.globl hello
.type hello, @function
hello:
.LFB2:
pushq %rbp
.LCFI0:
movq %rsp, %rbp
.LCFI1:
subq $16, %rsp
.LCFI2:
movq %rdi, -8(%rbp)
movl %esi, -12(%rbp)
movl %edx, -16(%rbp)
movl -16(%rbp), %ecx
movl -12(%rbp), %edx
movq -8(%rbp), %rsi
movl $.LC0, %edi
movl $0, %eax
```

```
call printf
leave
ret
.LFE2:
.size hello, .-hello
<..snip..>
```

Registers,
constants, &
labels

What are Pseudo-Instructions?

Assembler directives, with various purposes

Data & instruction encoding:

- ♦ Separate instructions & data into sections
- ♦ Reserve memory with initial data values
- ♦ Reserve memory w/o initial data values
- ♦ Align instructions & data

Provide information useful to linker or debugger

- ♦ Correlate source code with assembly/machine
- ♦ ...

Instructions & Pseudo-Instructions

```
.file    "main.c"
.section .rodata

.LC0:
.string "Hello, %s, it's %d:%02d."
.text
.globl hello
.type    hello, @function

hello:
.LFB2:
    pushq    %rbp

.LCFI0:
    movq     %rsp, %rbp

.LCFI1:
    subq     $16, %rsp

.LCFI2:
    movq     %rdi, -8(%rbp)
    movl     %esi, -12(%rbp)
    movl     %edx, -16(%rbp)
    movl     -16(%rbp), %ecx
    movl     -12(%rbp), %edx
    movq     -8(%rbp), %rsi
    movl     $.LC0, %edi
    movl     $0, %eax
```

```
    call     printf
    leave
    ret

.LFE2:
    .size    hello, .-hello
<..snip..>
```

Instructions,
Pseudo-Instructions,
& Label Definitions

Label Types

```
.file    "main.c"
.section .rodata

.LC0:
.string "Hello, %s, it's %d:%02d."
.text
.globl hello
.type    hello, @function

hello:
.LFB2:
    pushq    %rbp
.LCFI0:
    movq     %rsp, %rbp
.LCFI1:
    subq     $16, %rsp
.LCFI2:
    movq     %rdi, -8(%rbp)
    movl     %esi, -12(%rbp)
    movl     %edx, -16(%rbp)
    movl     -16(%rbp), %ecx
    movl     -12(%rbp), %edx
    movq     -8(%rbp), %rsi
    movl     $.LC0, %edi
    movl     $0, %eax
```

```
    call     printf
    leave
    ret
.LFE2:
    .size    hello, .-hello
<..snip..>
```

Definitions,
internal references,
& external references

Assembly/Machine Language – Semantics

Basic model of execution

- ♦ **Fetch instruction, from memory @ PC**
- ♦ **Increment PC**
- ♦ **Decode instruction**
- ♦ **Fetch operands, from registers or memory**
- ♦ **Execute operation**
- ♦ **Store result(s), in registers or memory**

Recall C Program

```
#include <stdio.h>

void
hello(char *name, int hour, int min)
{
    printf("Hello, %s, it's %d:%02d.",
           name, hour, min);
}

int
main(void)
{
    hello("Alan", 2, 55);

    return (0);
}
```

Program Execution

```
.file    "main.c"
.section .rodata
.LC0:
.string "Hello, %s, it's %d:%02d."
.text
.globl hello
.type    hello, @function
hello:
.LFB2:
        pushq   %rbp
.LCFI0:
        movq    %rsp, %rbp
.LCFI1:
        subq    $16, %rsp
.LCFI2:
        movq    %rdi, -8(%rbp)
        movl    %esi, -12(%rbp)
        movl    %edx, -16(%rbp)
        movl    -16(%rbp), %ecx
        movl    -12(%rbp), %edx
        movq    -8(%rbp), %rsi
        movl    $.LC0, %edi
        movl    $0, %eax
```

```
        call    printf
        leave   %rsp
        ret
.LFE2:
        .size   hello, .-hello
        .section .rodata
.LC1:
        .string "Alan"
        .text
.globl main
.type    main, @function
main:
.LFB3:
        pushq   %rbp
.LCFI3:
        movq    %rsp, %rbp
.LCFI4:
        movl    $55, %edx
        movl    $2, %esi
        movl    $.LC1, %edi
        call    hello
        movl    $0, %eax
<..next slide..>
```

Program Execution (cont.)

```
        movl    $0, %eax
        leave
        ret
.LFE3:
        .size   main, .-main
<..snip..>
```

More x86-64 Assembly

x86-64 notes on course web page

- ♦ **Simple assembly language manual**
- ♦ **Some example code**
- ♦ **Along with lecture notes, should cover everything you need for the course**

Some code sequences generated by the compiler can still be confusing

- ♦ **Usually not important for this class (web is helpful if you are still curious)**

Next Time

Program Linking