

Linking

Alan L. Cox
alc@cs.rice.edu

x86-64 Assembly

Brief overview last time

- ♦ Lecture notes and x86-64 assembly overview available on course web page

You will not write any x86-64 assembly

- ♦ Need to be able to recognize/understand code fragments
- ♦ Need to be able to correlate assembly code to source code

More assembly examples today

Linking

Linking: collecting and combining various pieces of code and data into a single file that can be loaded into memory and executed

Why learn about linking?

- ♦ **It will help you build large programs**
- ♦ **It will help you avoid dangerous program errors**
- ♦ **It will help you understand how language scoping rules are implemented**
- ♦ **It will help you understand other important systems concepts (that are covered later in the class)**
- ♦ **It will enable you to exploit shared libraries**

Example Program

```
/* main.c */
void swap(void);
int buf[2] = {1, 2};

int main(void)
{
    swap();
    return (0);
}
```

```
/* swap.c */
extern int buf[];

int *bufp0 = &buf[0];
int *bufp1;

void swap(void)
{
    int temp;

    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}
```

Compilation

Compiler: c to o

UNIX% gcc -v -O -g -o p main.c swap.c

cc1 -quiet -v main.c -quiet -dumpbase main.c -mtune=generic
-auxbase main -g -O -version -o /tmp/cchnheja.s

as -V -Qy -o /tmp/ccmNFRZd.o /tmp/cchnheja.s

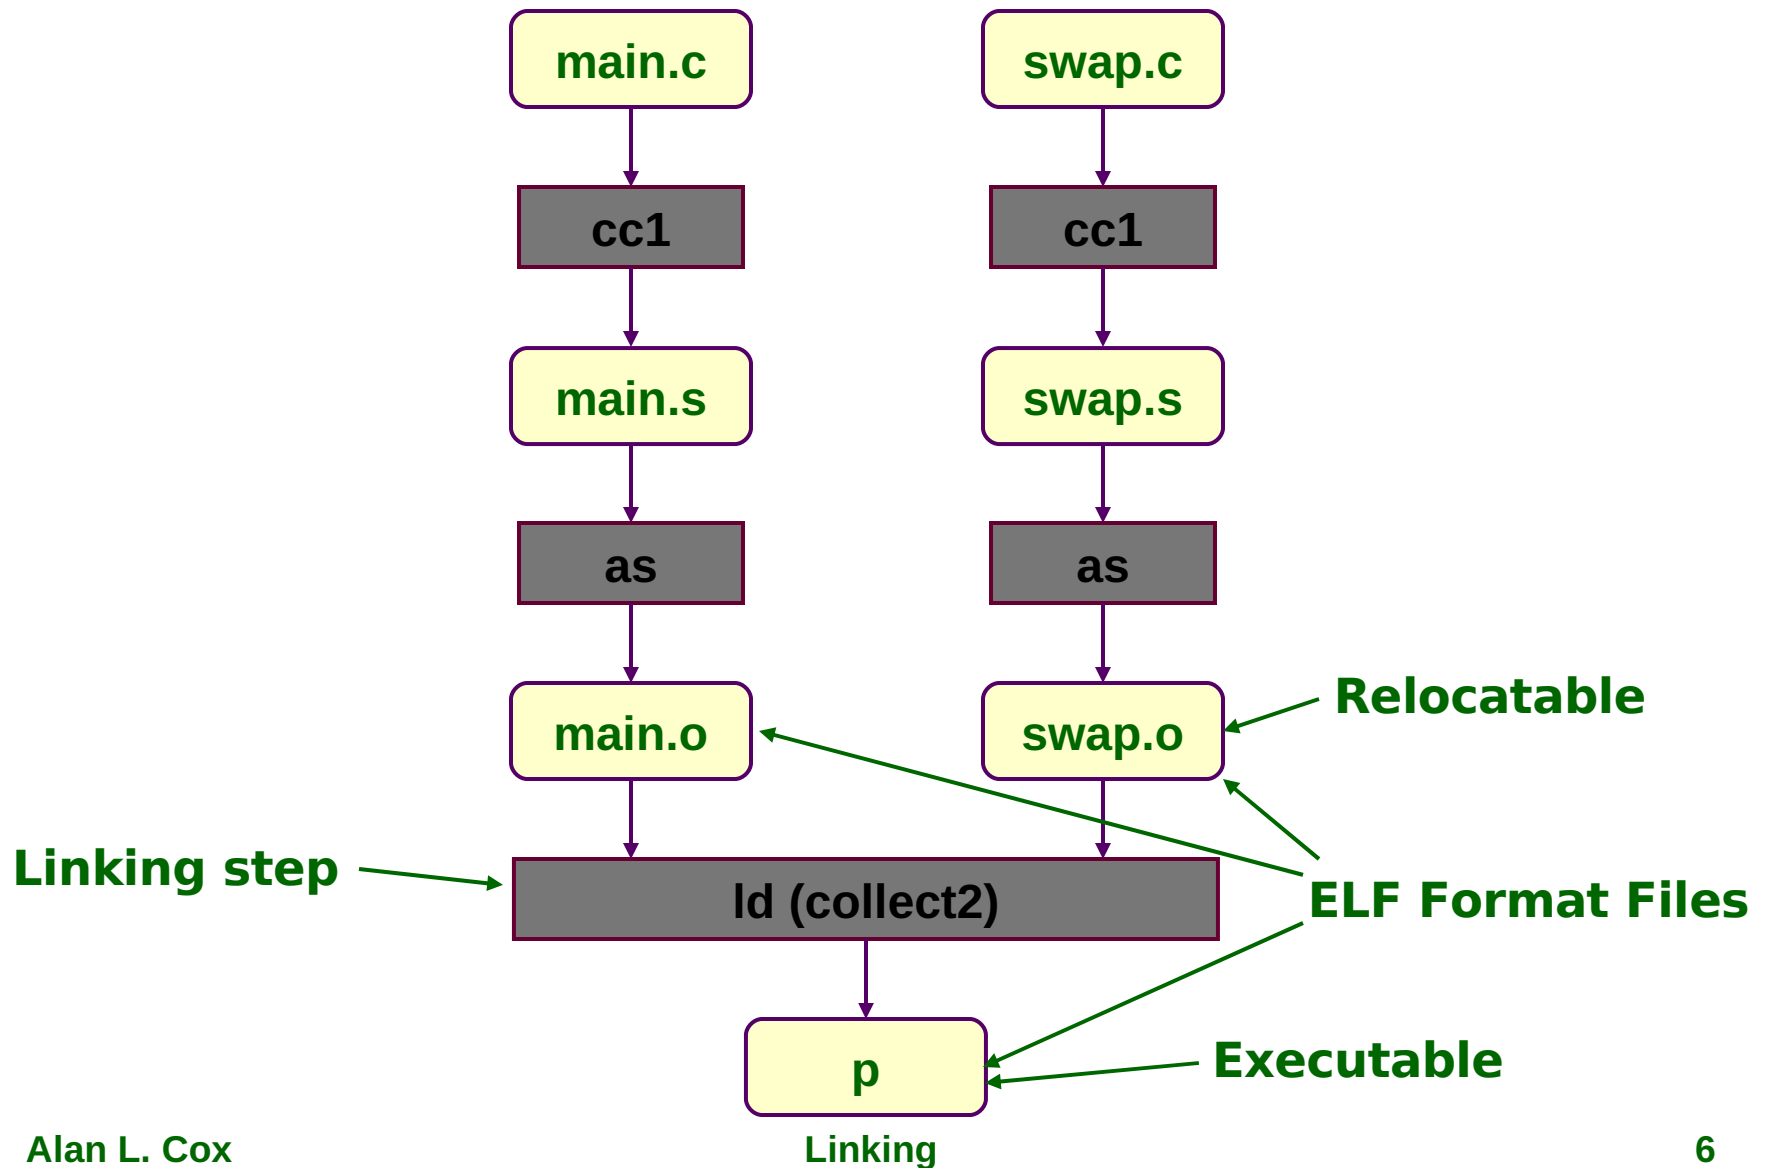
cc1 -quiet -v swap.c -quiet -dumpbase swap.c -mtune=generic
-auxbase swap -g -O -version -o /tmp/cchnheja.s

as -V -Qy -o /tmp/ccx8FECg.o /tmp/ccheheja.s

collect2 --eh-frame-hdr -m elf_x86_64 --hash-style=gnu -dynamic-
linker /lib64/ld-linux-x86-64.so.2 -o p crt1.o crti.o
crtbegin.o -L<..snip..> /tmp/ccmNFRZd.o /tmp/ccx8FECg.o -lgcc
--as-needed -lgcc_s --no-as-needed -lc -lgcc --as-needed
-lgcc_s --no-as-needed crtend.o crtn.o

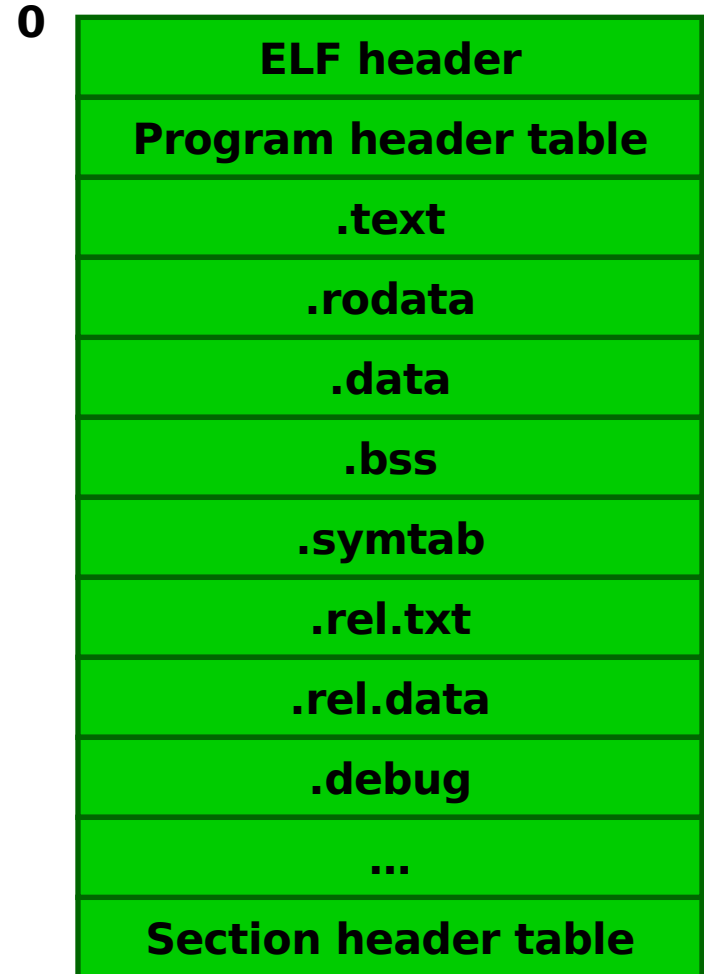
Linker: .o to executable

Compilation



ELF File Format

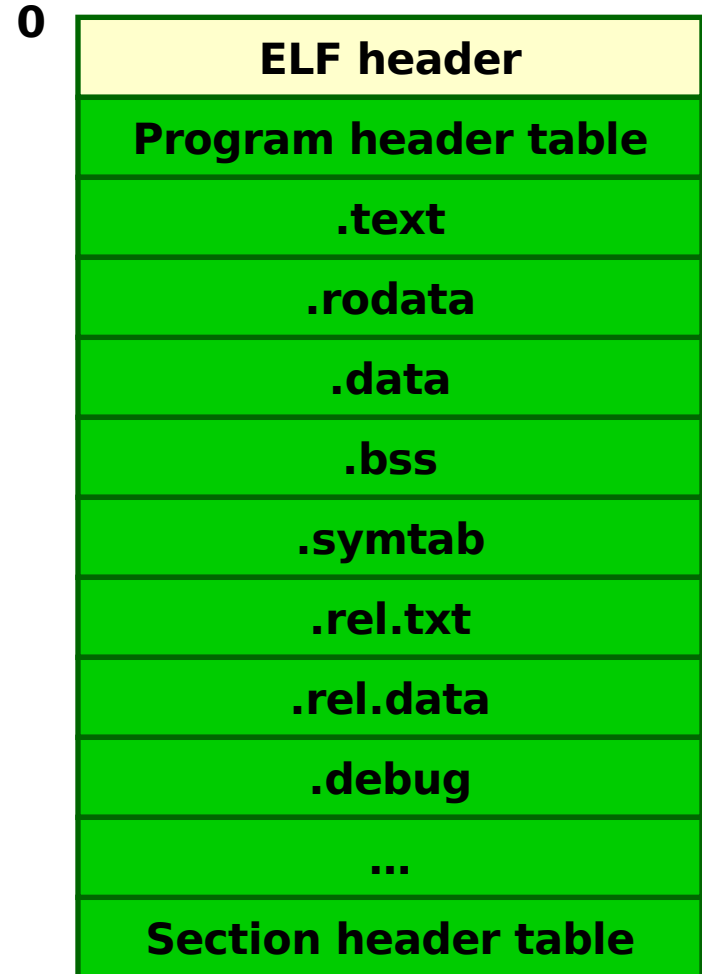
Order & existence of segments is arbitrary, except ELF header must be present and first



ELF Header

Basic description of file contents:

- ♦ File format identifier
- ♦ Endianness
- ♦ Alignment requirement for other sections
- ♦ Location of other sections
- ♦ Code's starting address
- ♦ ...



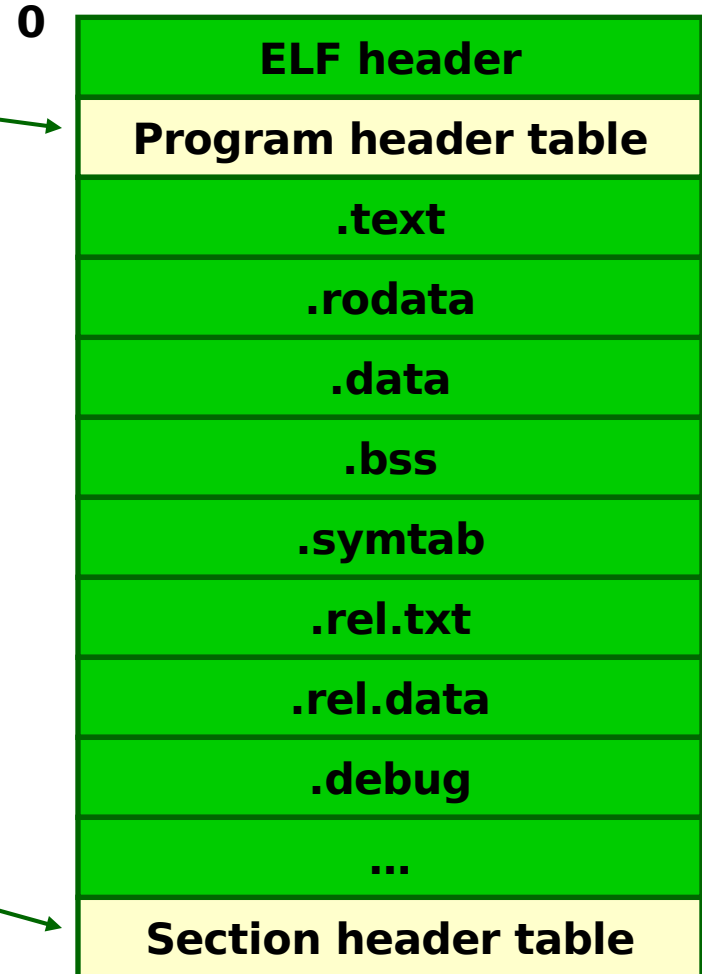
Program and Section Headers

**Info about other sections
necessary for loading**

- ♦ Required for executables & libraries

**Info about other sections
necessary for linking**

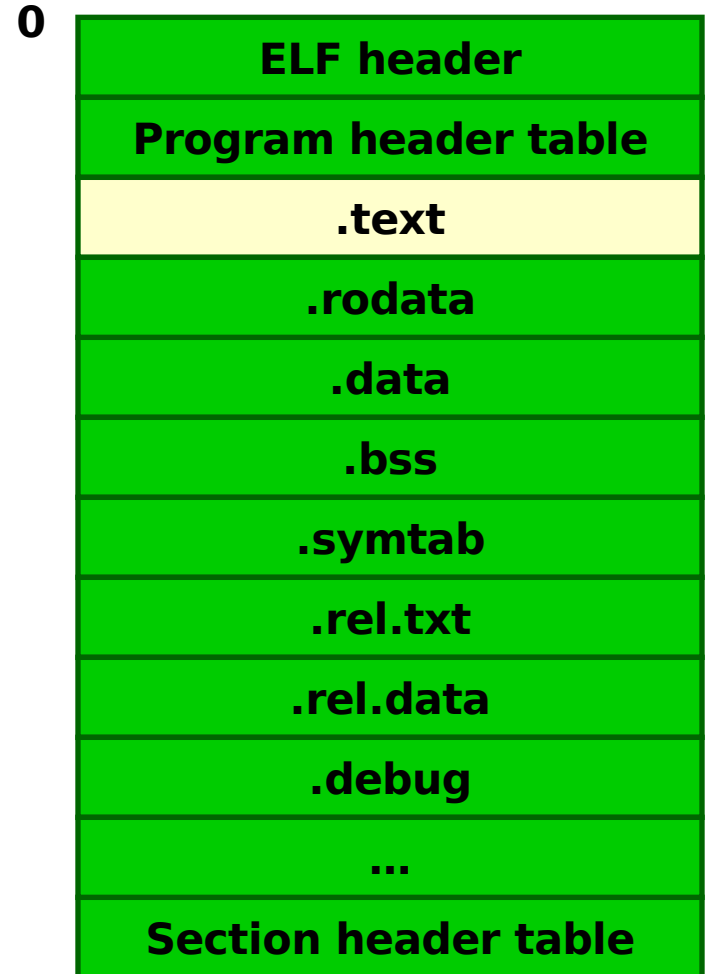
- ♦ Required for relocatables



Text Section

Code

- ♦ read-only



Data Sections

Static data

- ♦ initialized, read-only
- ♦ initialized, read/write
- ♦ uninitialized, read/write (BSS = “Block Started by Symbol” pseudo-op for IBM 704)

Initialized

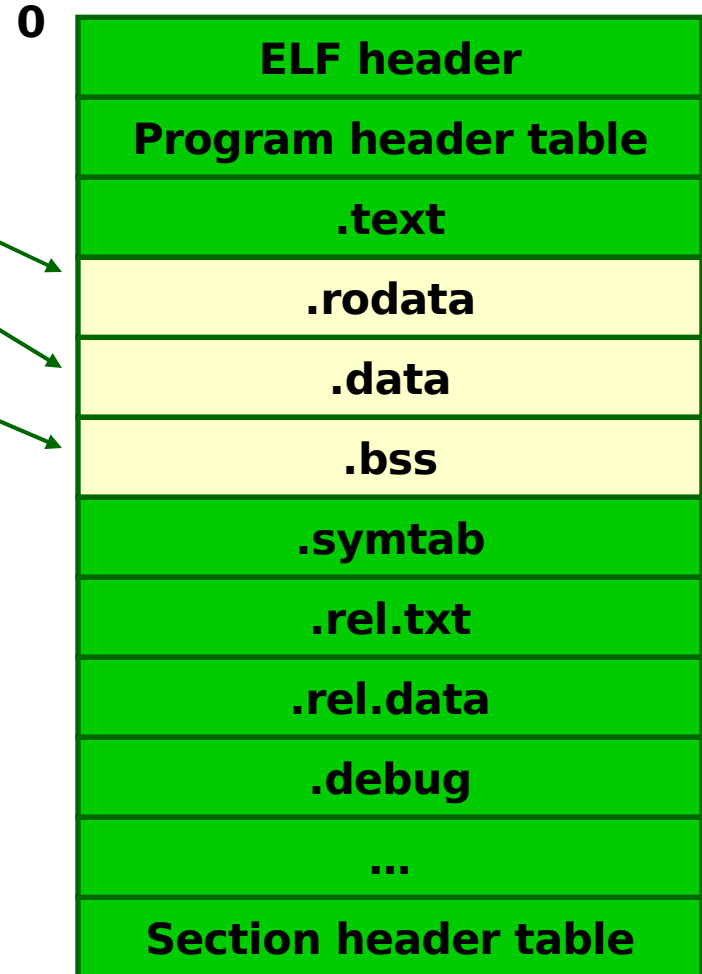
- ♦ Initial values in ELF file

Uninitialized

- ♦ Only total size in ELF file

Writable distinction enforced at run-time

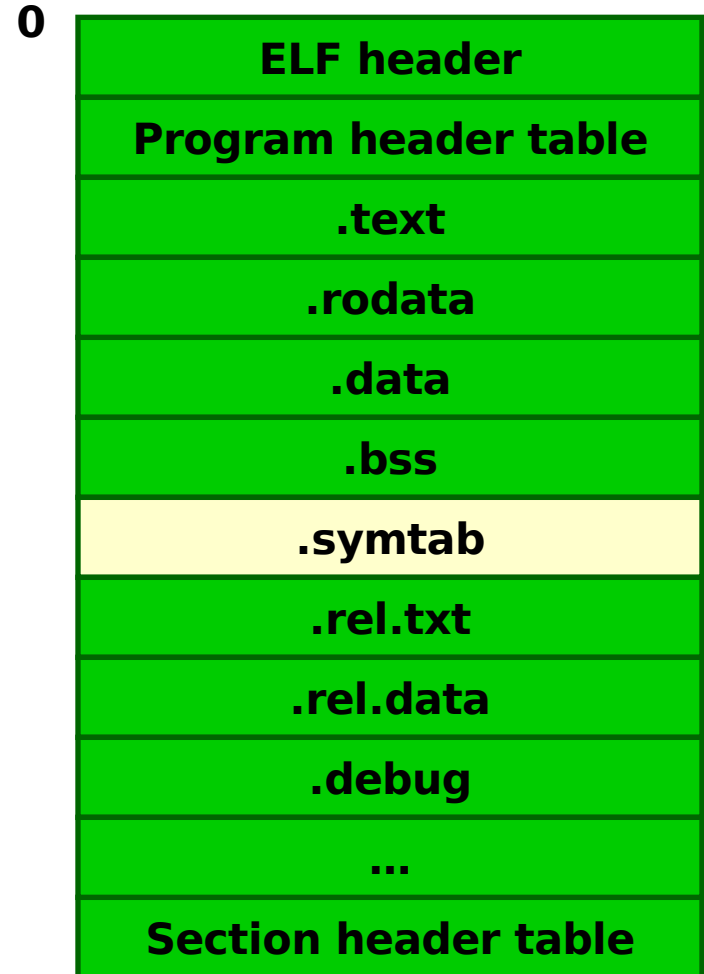
- ♦ Why? Protection; sharing
- ♦ How? Virtual memory



Symbol Table

Describes where global variables and functions are defined

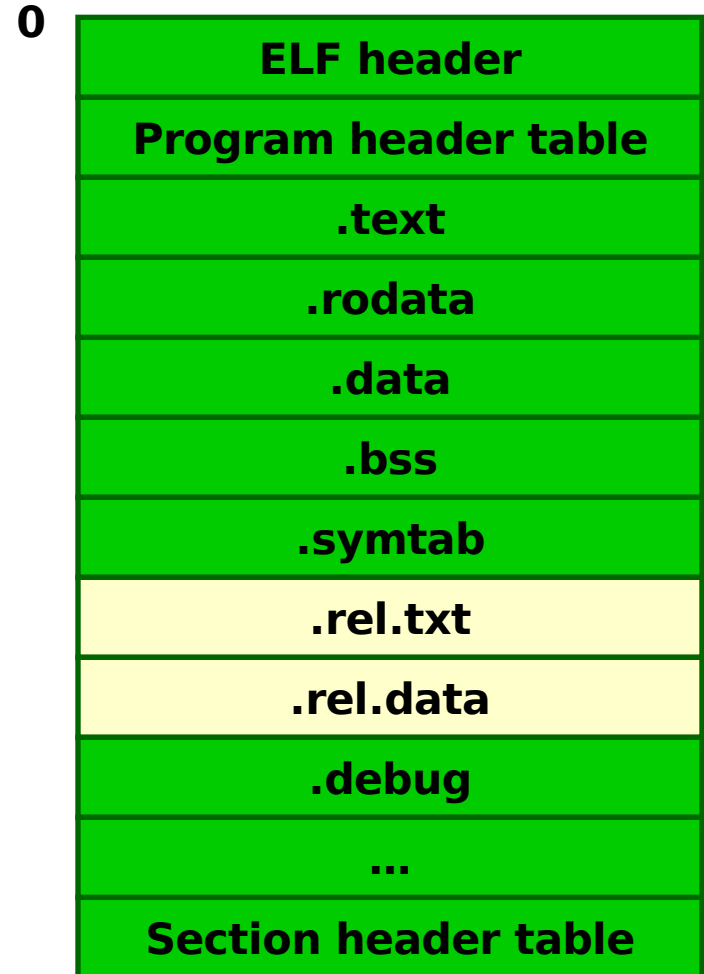
- ♦ **Present in all relocatable ELF files (not just in files compiled for debugging)**



Relocation Information

Describes where and how labels are used

- ♦ **Allows object files to be easily relocated**



Debug Section

Relates source code to the object code within the ELF file

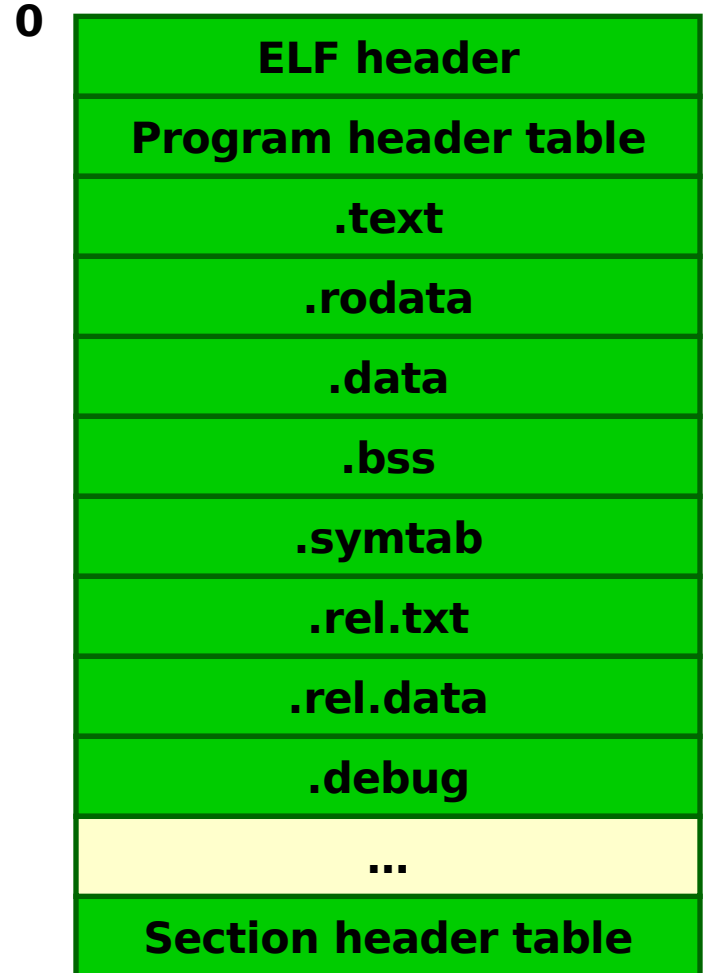
0

ELF header
Program header table
.text
.rodata
.data
.bss
.symtab
.rel.txt
.rel.data
.debug
...
Section header table

Other Sections

Other kinds of sections also supported, including:

- ♦ **Other debugging info**
- ♦ **Version control info**
- ♦ **Dynamic linking info**
- ♦ **C++ initializing & finalizing code**



Linker Symbol Classification

Global symbols

- ♦ Symbols defined by module *m* that can be referenced by other modules
- ♦ C: non-static functions & global variables

External symbols

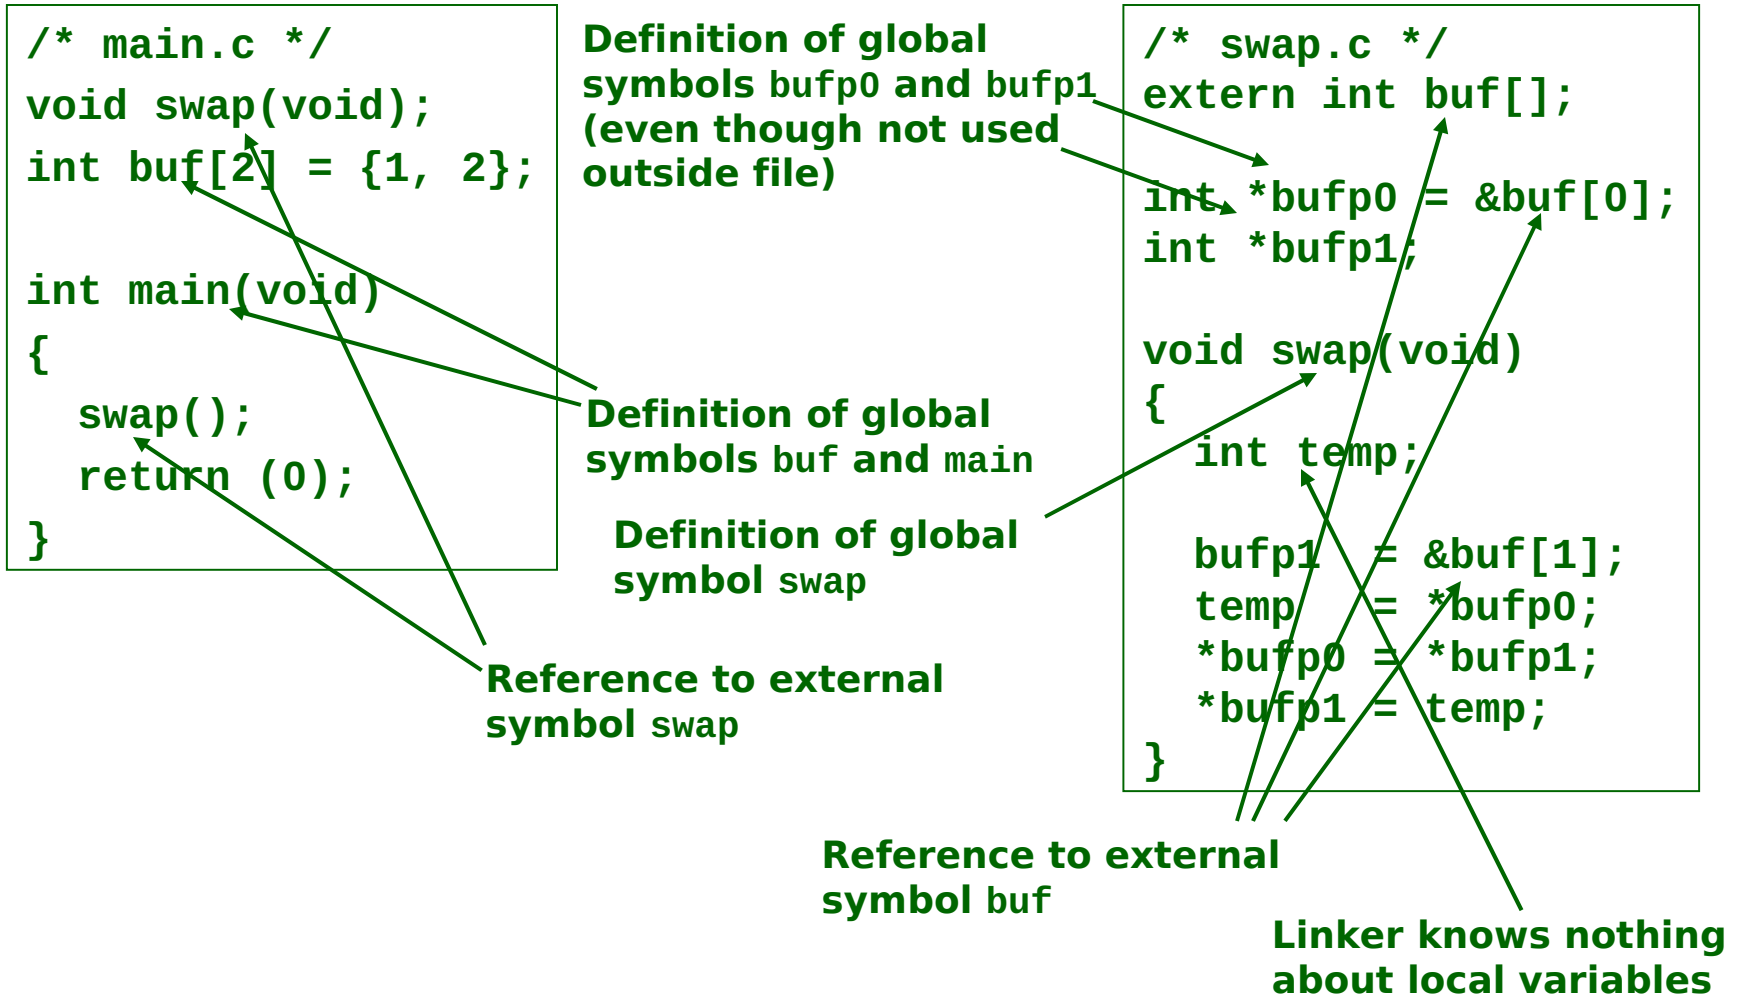
- ♦ Symbols referenced by module *m* but defined by some other module
- ♦ C: extern functions & variables

Local symbols

- ♦ Symbols that are defined and referenced exclusively by module *m*
- ♦ C: static functions & variables

Local linker symbols \neq local program variables!

Linker Symbols



Linking: Symbols

```
/* main.c */
void swap(void);
int buf[2] = {1, 2};

int main(void)
{
    swap();
    return (0);
}
```

What's missing?

- ♦ **swap** – where is it?

swap is a function that is not defined in the object file
it is defined in the library (libc) and is set 0
offset of 1 (NDX)1 (.text)

use **readelf -S** to see sections

```
UNIX% gcc -O -c main.c
UNIX% readelf -s main.o
```

Symbol table '.symtab' contains 11 entries:

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
...							
8:	0000000000000000	19	FUNC	GLOBAL	DEFAULT	1	main
9:	0000000000000000	0	NOTYPE	GLOBAL	DEFAULT	UND	swap
10:	0000000000000000	8	OBJECT	GLOBAL	DEFAULT	3	buf

Linking: Symbols

```
/* swap.c */
extern int buf[];
int *bufp0 = &buf[0];
int *bufp1;

void swap(void)
{
    int temp;

    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}
```

What's missing?

- ♦ **buf** – where is it?

Unaligned (32-bit) alignment requirement

Symbol table '.symtab' contains 12 entries:

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
8:	000000000000000000	38	FUNC	GLOBAL	DEFAULT	1	swap
9:	000000000000000000	0	NOTYPE	GLOBAL	DEFAULT	UND	buf
10:	000000000000000008	8	OBJECT	GLOBAL	DEFAULT	COM	bufp1
11:	000000000000000000	8	OBJECT	GLOBAL	DEFAULT	3	bufp0

Name Mangling

Other languages (i.e. Java and C++) allow overloaded methods

- ♦ Functions then have the same name but take different numbers/types of arguments
- ♦ How does the linker disambiguate these symbols?

Generate unique names through *mangling*

- ♦ Mangled names are compiler dependent
- ♦ Example: class “Foo”, method “bar(int, long)”:
 - bar__3Fooil
 - _ZN3Foo3BarEil
- ♦ Similar schemes are used for global variables, etc.

Linking Steps

Symbol Resolution

- ♦ **Determine where symbols are located and what size data/code they refer to**

Relocation

- ♦ **Combine modules, relocate code/data, and fix symbol references based on new locations**


Symbol Resolution

Undefined symbols must be resolved

- ♦ Where are they located
- ♦ What size are they?

Linker looks in the symbol tables of all relocatable object files

- ♦ Assuming every unknown symbol is defined once and only once, this works well



Linker Relocation

Once all symbols are resolved, must combine the input files

- ♦ **Total code size is known**
- ♦ **Total data size is known**
- ♦ **All symbols must be assigned run-time addresses**


Sections must be merged

- ♦ **Only one text, data, etc. section in final executable**
- ♦ **Final run-time addresses of all symbols are defined**

Symbol references must be corrected

- ♦ **All symbol references must now refer to their actual locations**

Relocation: Merging Files



Linking: Relocation

```
/* main.c */  
void swap(void);  
int buf[2] = {1, 2};  
  
int main(void)  
{  
    swap();  
    return (0);  
}
```

can also use `readelf
-r` to see relocation
information

```
UNIX% objdump -r -d main.o
```

```
main.o:      file format elf64-x86-64
```

```
Disassembly of section .text:
```

```
0000000000000000 <main>:
```

0:	48 83 ec 08	sub	\$0x8,%rsp
4:	e8 00 00 00 00	callq	9 <main+0x9>
5:		R_X86_64_PC32	
		swap+0xfffffffffffffc	
9:	b8 10 20 00	mov	\$0x0,%eax
e:	48 83 c4 08	add	\$0x8,%rsp
12:	c3	retq	

Type of instruction (R_X86_64_PC32) and
Offset in text section (relative to the start of the section)
is stored in a different section of the file)

Linking: Relocation

```
/* swap.c */
extern int buf[];
int *bufp0 = &buf[0];
int *bufp1;

void swap()
{
    int temp;

    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}
```

```
UNIX% objdump -r -D swap.o
```

```
swap.o:      file format elf64-x86-64
```

```
Disassembly of section .text:
```

```
0000000000000000 <swap>:
```

```
0:  48 c7 05 00 00 00 00 movq $0x0,0(%rip)
```

```
7:  00 00 00 00
```

```
3: R_X86_64_PC32
```

```
bufp1+0xfffffffffffffffff8
```

```
7: R_X86_64_32S buf+0x4
```

```
<..snip..>
```

```
Disassembly of section .data:
```

```
0000000000000000 <bufp0>:
```

```
...
```

```
0: R_X86_64_64 buf
```

Need relocated address of bufp1 to initialize bufp0 with &buf[0] (== buf)

After Relocation

0000000000000000 <main>:

```
0:  48 83 ec 08      sub    $0x8,%rsp
4:  e8 00 00 00 00    callq 9 <main+0x9>
                    5: R_X86_64_PC32 swap+0xfffffffffffffc
9:  b8 10 20 00      mov    $0x0,%eax0
e:  48 83 c4 08      add    $0x8,%rsp
12: c3               retq
```



0000000000400448 <main>:

```
400448:  48 83 ec 08      sub    $0x8,%rsp
40044c:  e8 0b 00 00 00    callq 40045c <swap>
400451:  b8 10 20 00      mov    $0x0,%eax0
400456:  48 83 c4 08      add    $0x8,%rsp
40045a:  c3               retq
```

After Relocation

```
0000000000000000 <swap>:
 0:  48 c7 05 00 00 00 00 00 movq $0x0,0(%rip)
 7:  00 00 00 00

                3: R_X86_64_PC32 bufp1+0xfffffffffffffffff8
                7: R_X86_64_32S  buf+0x4

<..snip..>
0000000000000000 <bufp0>:
    ...
                0: R_X86_64_64 buf
```



```
000000000040045c <swap>:
40045c:  48 c7 05 01 04 20 00 00 movq $0x600848,2098177(%rip)
400463:  48 08 60 00                                # 600868 <bufp1>
<..snip..>
0000000000600850 <bufp0>:
600850:  44 08 60 00 00 00 00 00
```

Problem: Undefined Symbols

```
UNIX% gcc -O2 -o p main.c
/tmp/cccpTy0d.o: In function `main':
main.c:(.text+0x5): undefined reference to `swap'
collect2: ld returned 1 exit status
UNIX%
```

Missing symbols are not compiler errors

- ♦ May be defined in another file
- ♦ Compiler just inserts an undefined entry in the symbol table

During linking, any undefined symbols that cannot be resolved cause an error

Problem: Multiply Defined Symbols

Different files could define the same symbol

- ♦ **Is this an error?**
- ♦ **If not, which one should be used? One or many?**

Linking: Example

```
int x = 3;  
int y = 4;  
int z;  
  
int foo(int a) {...}  
int bar(int b) {...}
```



```
extern int x;  
static int y = 6;  
int z;  
  
int foo(int a);  
static int bar(int b) {...}
```



?

?

Note: Linking uses object files
Examples use source-level for
convenience

Linking: Example

```
int x = 3;  
int y = 4;  
int z;  
  
int foo(int a) {...}  
int bar(int b) {...}
```

Defined in one file



```
extern int x;  
static int y = 6;  
int z;  
  
int foo(int a);  
static int bar(int b) {...}
```

Declared in other files



```
int x = 3;  
  
  
  
  
  
  
  
  
  
int foo(int a) {...}
```

Only one copy exists

Linking: Example

```
int x = 3;  
int y = 4;  
int z;  
  
int foo(int a) {...}  
int bar(int b) {...}
```



```
extern int x;  
static int y = 6;  
int z;  
  
int foo(int a);  
static int bar(int b) {...}
```



```
int x = 3;  
int y = 4;  
int y' = 6;  
  
int foo(int a) {...}  
int bar(int b) {...}  
int bar'(int b) {...}
```

Private names
not in symbol
table.
Can't conflict
with other files'
names

Renaming is a
convenient source-
level way to
understand this

Linking: Example

```
int x = 3;  
int y = 4;  
int z;  
  
int foo(int a) {...}  
int bar(int b) {...}
```



```
extern int x;  
static int y = 6;  
int z;  
  
int foo(int a);  
static int bar(int b) {...}
```



```
int x = 3;  
int y = 4;  
int y' = 6;  
int z;  
  
int foo(int a) {...}  
int bar(int b) {...}  
int bar'(int b) {...}
```

C allows you to
omit “**extern**” in
some cases –
Don't!

Strong & Weak Symbols

Program symbols are either strong or weak

<i>strong</i>	procedures & initialized globals
<i>weak</i>	uninitialized globals

p1.c

strong	→	int foo=5;
strong	→	p1() {}

p2.c

int foo;	←	weak
p2() {}	←	strong

Strong & Weak Symbols

A strong symbol can only appear once

A weak symbol can be overridden by a strong symbol of the same name

- ♦ **References to the weak symbol resolve to the strong symbol**

If there are multiple weak symbols, the linker can pick an arbitrary one

Linker Puzzles: What Happens?

```
int x;  
p1() {}
```

```
p1() {}
```

Link time error: two strong symbols **p1**

```
int x;  
p1() {}
```

```
int x;  
p2() {}
```

References to **x** will refer to the same uninitialized int.

Is this what you really want?

```
int x;  
int y;  
p1() {}
```

```
double x;  
p2() {}
```

Writes to **x** in **p2** might overwrite **y**!
Evil!

```
int x=7;  
int y=5;  
p1() {}
```

```
double x;  
p2() {}
```

Writes to **x** in **p2** will overwrite **y**!
Nasty!

```
int x=7;  
p1() {}
```

```
int x;  
p2() {}
```

References to **x** will refer to the same initialized variable

Nightmare scenario: replace **int** with a **struct** type, compile each file with different alignment

rules
Linking

Libraries

How should functions commonly used by programmers be provided?

- ♦ **Math, I/O, memory management, string manipulation, etc.**
- ♦ **Option 1: Put all functions in a single source file**
 - Programmers link big object file into their programs
 - Space and time inefficient
- ♦ **Option 2: Put each function in a separate source file**
 - Programmers explicitly link appropriate object files into their programs
 - More efficient, but burdensome on the programmer

Solution: static libraries (.a archive files)

- ♦ **Multiple relocatable files + index → single archive file**
- ♦ **Only links the subset of relocatable files from the library that are used in the program**
- ♦ **Example: `gcc -o fpmath main.c float.c -lm`**

Two Common Libraries

`libc.a` (the C standard library)

- ♦ 4 MB archive of 1395 object files
- ♦ I/O, memory allocation, signal handling, string handling, data and time, random numbers, integer math
- ♦ Usually automatically linked

`libm.a` (the C math library)

- ♦ 1.3 MB archive of 401 object files
- ♦ floating point math (sin, cos, tan, log, exp, sqrt, ...)
- ♦ Use “-lm” to link with your program

```
UNIX% ar t /usr/lib64/libc.a
```

```
...  
fprintf.o  
...  
feof.o  
...  
fputc.o  
...  
strlen.o  
...
```

```
UNIX% ar t /usr/lib64/libm.a
```

```
...  
e_sinh.o  
e_sqrt.o  
e_gamma_r.o  
k_cos.o  
k_rem_pio2.o  
k_sin.o  
k_tan.o  
...
```

Creating a Library

```
/* vector.h */  
void addvec(int *x, int *y, int *z, int n);  
void multvec(int *x, int *y, int *z, int n);
```

```
/* addvec.c */  
#include "vector.h"  
void addvec(int *x, int *y,  
            int *z, int n)  
{  
    int i;  
  
    for (i = 0; i < n; i++)  
        z[i] = x[i] + y[i];  
}
```

```
/* multvec.c */  
#include "vector.h"  
void multvec(int *x, int *y,  
             int *z, int n)  
{  
    int i;  
  
    for (i = 0; i < n; i++)  
        z[i] = x[i] * y[i];  
}
```

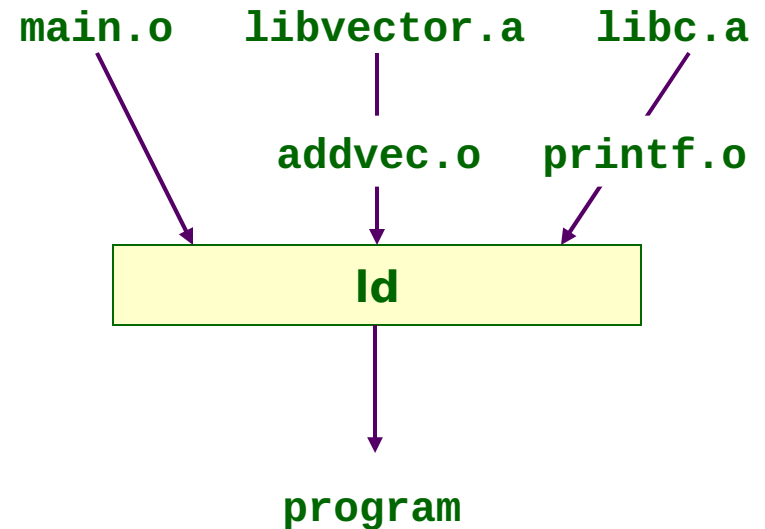
```
UNIX% gcc -c addvec.c multvec.c  
UNIX% ar rcs libvector.a addvec.o multvec.o
```


Using a library

```
/* main.c */
#include <stdio.h>
#include "vector.h"

int x[2] = {1, 2};
int y[2] = {3, 4};
int z[2];

int main(void)
{
    addvec(x, y, z, 2);
    printf("z = [%d %d]\n", z[0], z[1]);
    return (0);
}
```



```
UNIX% gcc -O2 -c main.c
UNIX% gcc -static -o program main.o ./libvector.a
```

How to Link: Basic Algorithm

Keep a list of the current unresolved references.

For each object file (.o and .a) in command-line order

- ♦ Try to resolve each unresolved reference in list to objects defined in current file
- ♦ Try to resolve each unresolved reference in current file to objects defined in previous files
- ♦ Concatenate like sections (.text with .text, etc.)

If list empty, output executable file, else error

Problem: Command line order matters! Link libraries last:

```
UNIX% gcc main.o libvector.a
UNIX% gcc libvector.a main.o
main.o: In function `main':
main.o(.text+0x4): undefined reference to `addvec'
```

Dynamic Libraries

Static

Linked at compile-time

UNIX: foo.a

Relocatable ELF File

Dynamic

Linked at run-time

UNIX: foo.so

Shared ELF File

What are the differences?

Static & Dynamic Libraries

Static

- ♦ Library code added to executable file
- ♦ Larger executables
- ♦ Must recompile to use newer libraries
- ♦ Executable is self-contained
- ♦ Some time to load libraries at compile-time
- ♦ Library code shared only among copies of same program

Dynamic

- ♦ Library code not added to executable file
- ♦ Smaller executables
- ♦ Uses newest (or smallest, fastest, ...) library without recompiling
- ♦ Depends on libraries at run-time
- ♦ Some time to load libraries at run-time
- ♦ Library code shared among all uses of library

Static & Dynamic Libraries

Static

Creation

```
ar r libfoo.a bar.o baz.o
ranlib libfoo.a
```

Use

```
gcc -o zap zap.o -lfoo
```

Adds library's code, data,
symbol table, relocation
info, ...

Dynamic

Creation

```
gcc -shared -Wl,-soname,libfoo.so  
-o libfoo.so bar.o baz.o
```

Use

```
gcc -o zap zap.o -lfoo
```

Adds library's symbol table,
relocation info

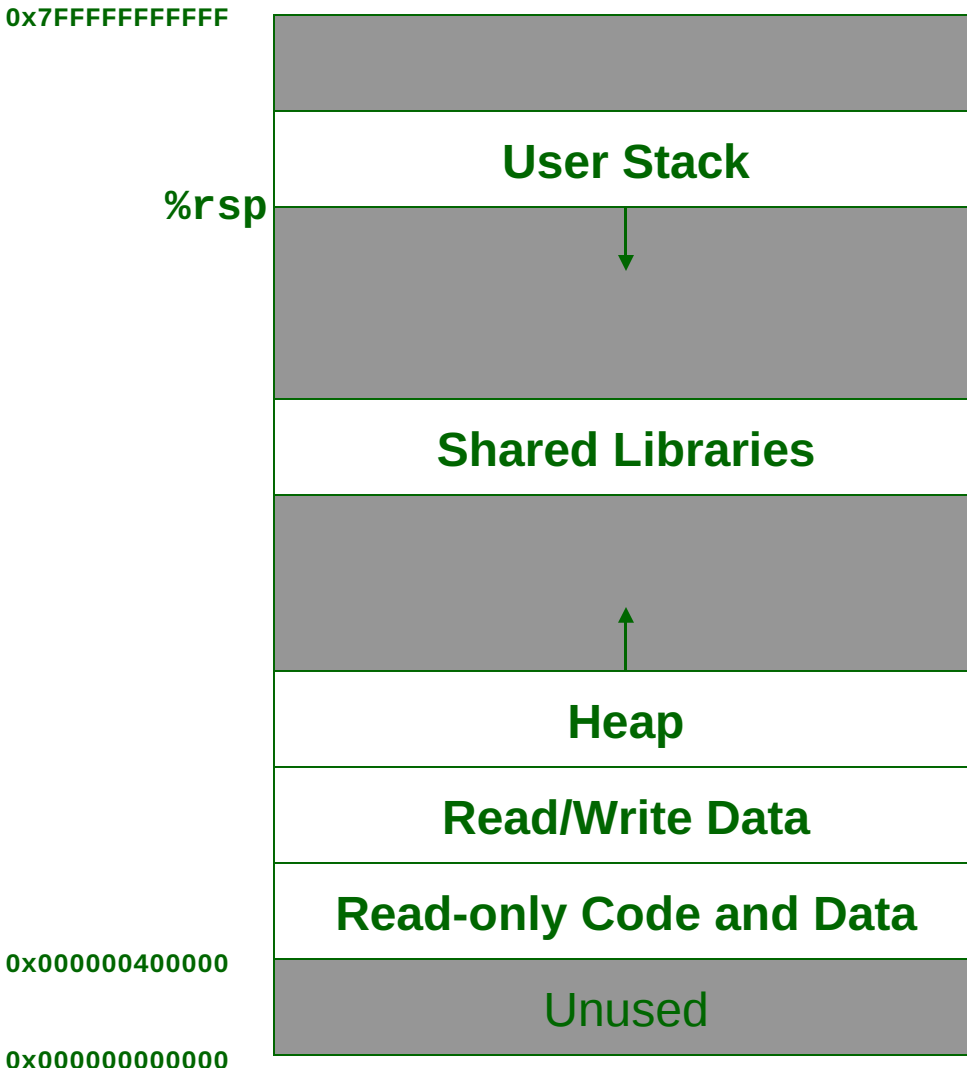
Loading

Linking yields an executable that can actually be run

Running a program

- ♦ **unix% ./program**
- ♦ **Shell does not recognize “program” as a shell command, so assumes it is an executable**
- ♦ **Invokes the *loader* to load the executable into memory (any unix program can invoke the loader with the *execve* function – more later)**

Creating the Memory Image (sort of...)



Create code and data segments

- ♦ Copy code and data from executable into these segments

Create initial heap segment

- ♦ Grows up from read/write data

Create stack

- ♦ Starts near the top and grows downward

Call dynamic linker to load shared libraries and relocate references

Starting the Program

Jump to program's entry point (stored in ELF header)

- ♦ For C programs, this is the `_start` symbol

Execute `_start` code (from `crt1.o` – same for all C programs)

- ♦ call `__libc_init_first`
- ♦ call `_init`
- ♦ call `atexit`
- ♦ call `main`
- ♦ call `_exit`

Position Independent Code

Static libraries compile with unresolved global & local addresses

- ♦ **Library code & data concatenated & addresses resolved when linking**

Position Independent Code

By default (in C), dynamic libraries compile with resolved global & local addresses

- ♦ **E.g., `libfoo.so` starts at `0x400000` in every application using it**
- ♦ **Advantage: Simplifies sharing**
- ♦ **Disadvantage: Inflexible – must decide ahead of time where each library goes, otherwise libraries can conflict**

Position Independent Code

Can compile dynamic libraries with unresolved global & local addresses

- ♦ `gcc -shared -fPIC ...`
- ♦ **Advantage: More flexible – no conflicts**
- ♦ **Disadvantage: Code less efficient – referencing these addresses involves indirection**

Library Interpositioning

Linking with non-standard libraries that use standard library symbols

- ♦ **“Intercept” calls to library functions**

Some applications:

- ♦ **Security**
 - **Confinement (sandboxing)**
 - **Behind the scenes encryption**
 - **Automatically encrypt otherwise unencrypted network connections**
- ♦ **Monitoring & Profiling**
 - **Count number of calls to functions**
 - **Characterize call sites and arguments to functions**
 - **malloc tracing**
 - **Detecting memory leaks**
 - **Generating malloc traces**

Dynamic Linking at Run-Time

Application access to dynamic linker via API:

```
#include <dlfcn.h>

void
dlink(void)
{
    void *handle = dlopen("mylib.so", RTLD_LAZY);
    /* type */ myfunc = dlsym(handle, "myfunc");
    myfunc(...);
    dlclose(handle);
}
```

Symbols resolved
at first use, not
now

Error-checking omitted for clarity

Next Time

Exceptions