

# Measuring Time

---

**Alan L. Cox**  
**alc@cs.rice.edu**

Some slides adapted from CMU 15.213 slides


# Computer Time Scales

## Two Fundamental Time Scales

- ♦ **Processor:**  $\sim 10^{-9}$  s
- ♦ **External events:**  $\sim 10^{-2}$  s
  - Keyboard input
  - Disk seek
  - Screen refresh

## Implication

- ♦ Can execute many instructions while waiting for external event to occur
- ♦ Can alternate among processes without anyone noticing



# Measurement Challenge

---

## How Much Time Does Program X Require?

- ♦ **CPU time**
  - How many total seconds are used when executing X?
  - Measure used for most applications
  - Small dependence on other system activities
- ♦ **Actual (“Wall Clock”) Time**
  - How many seconds elapse between the start and the completion of X?
  - Depends on system load, I/O times, etc.

## Confounding Factors

- ♦ **How does time get measured?**
- ♦ **Many processes share computing resources**
  - Transient effects when switching from one process to another
  - Suddenly, the effects of alternating among processes become noticeable

# “Time” on a Computer System



real (wall clock) time

 = **user time** (*time executing instructions in the user process*)

 = **system time** (*time executing instructions in kernel on behalf of user process*)

 = **some other user's time** (*time executing instructions for a different user process*)


 +  +  = **real (wall clock) time**

*We will use the word “time” to refer to user time*


 **cumulative user time**

# Activity Periods: Light Load

- ♦ **Most of the time spent executing one process**
- ♦ **Periodic interrupts every 10ms**
  - Interval timer
  - Keep system from executing one process to exclusion of others
- ♦ **Other interrupts**
  - Due to I/O activity
- ♦ **Inactivity periods**
  - System time spent processing interrupts



# Activity Periods: Heavy Load



- ♦ **Sharing processor with one other active process**
- ♦ **From perspective of this process, system appears to be “inactive” for ~50% of the time**
  - **Other process is executing**

# Interval Counting

---


## OS Measures Runtimes Using Interval Timer

- ♦ **Maintain 2 counts per process**
  - User time
  - System time
- ♦ **Each timer interrupt, increment counter for executing process**
  - User time if running in user mode
  - System time if running in kernel mode

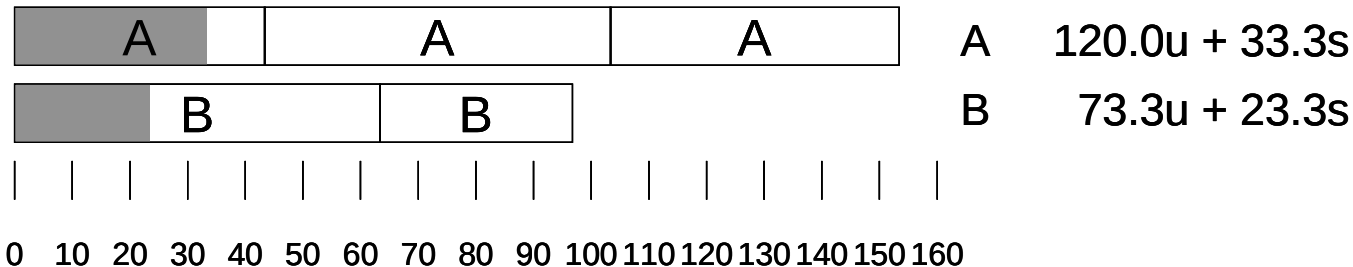
# Interval Counting Example

---

(a) Interval Timings



(b) Actual Times





# Unix time Command

---


```
unix% time make osevent
gcc -O2 -Wall -Wextra -g -c clock.c
gcc -O2 -Wall -Wextra -g -c options.c
gcc -O2 -Wall -Wextra -g -c load.c
gcc -O2 -Wall -Wextra -g -o osevent . . .
0.820u 0.300s 0:01.32 84.8%
```

- ♦ **0.82 seconds user time**
  - 82 timer intervals
- ♦ **0.30 seconds system time**
  - 30 timer intervals
- ♦ **1.32 seconds wall clock time**
- ♦ **84.8% of total was used running these processes**
  - $(.82+0.3)/1.32 = .848$

# Accuracy of Interval Counting

## Worst Case Analysis

- ♦ **Timer Interval =  $\delta$**
- ♦ **Single process segment measurement can be off by  $\pm \delta$**
- ♦ **No bound on error for multiple segments**
  - **Could consistently underestimate, or consistently overestimate**



Minimum

Maximum

- **Computed time = 70ms**


- **Min Actual =  $60 + \epsilon$**

- **Max Actual =  $80 - \epsilon$**

# Accuracy of Interval Counting

## Average Case Analysis

- ♦ Over/underestimates tend to balance out
- ♦ As long as total run time is sufficiently large
  - Min run time ~1 second
  - 100 timer intervals
- ♦ Consistently miss ~4% overhead due to timer interrupts



Minimum

Maximum

- Computed time = 70ms
- Min Actual =  $60 + \epsilon$
- Max Actual =  $80 - \epsilon$

# Cycle Counters

---

**Most modern systems have built in registers that are incremented every clock cycle**

- ♦ **Very fine grained**
- ♦ **Often counts elapsed global time**

**On x86 and x86-64 machines:**

- ♦ **64 bit counter**
  - **Cycle counter period:**
    - **A 3 GHz machine wraps around every 195 years**
- ♦ **Special instruction to access**

# x86 Cycle Counter

---

## RDTSC

- ◆ **Assembly instruction to access 64-bit cycle counter**
- ◆ **Places low/high 32 bits in two different registers**
- ◆ **Expressed as machine cycles, not nanoseconds**

```
uint64_t
rdtsc(void)
{
    uint32_t low, high;

    /* Get cycle counter */
    asm("rdtsc" : "=a" (low), "=d" (high)); /* %eax, %edx */
    return (low | ((uint64_t)high << 32));
}
```

# Measuring Cycles with rdtsc()

---

## Idea

- ♦ Get current cycle counter
- ♦ Compute something
- ♦ Get new cycle counter
- ♦ Perform 64-bit subtraction to get elapsed cycles

```
uint64_t start, end;
int i;
int iters = 100;

start = rdtsc();
for (i = 0; i < iters; i++)
    getpid();
end = rdtsc();

printf("getpid(): Average cycles = %ld\n", (end - start) / iters);
```

# Converting Cycles to Seconds

---

## Idea

- ♦ Compute elapsed cycles
- ♦ Get processor's clock frequency (cycles/second)
- ♦ Divide elapsed cycles by the clock frequency

## How do you get the clock frequency?

```
UNIX% cat /proc/cpuinfo
processor      : 0
vendor_id     : GenuineIntel
cpu family    : 6
model         : 23
model name    : Intel(R) Xeon(R) CPU           X5460   @ 3.16GHz
stepping      : 6
cpu MHz       : 3158.758
cache size    : 6144 KB
...
```

# Measurement Pitfalls

---

## Overhead/resolution

- ♦ Calling `rdtsc()` incurs some overhead
- ♦ Resolution of `rdtsc()` may not allow very short code sequences to be timed
- ♦ Want to measure long enough code sequence to compensate

## Unexpected Cache Effects

- ♦ artificial hits or misses
- ♦ e.g., these measurements were taken with the Alpha cycle counter:
  - ♦ `foo1(array1, array2, array3); /* 68,829 cycles */`
  - ♦ `foo2(array1, array2, array3); /* 23,337 cycles */`
- vs.
  - ♦ `foo2(array1, array2, array3); /* 70,513 cycles */`
  - ♦ `foo1(array1, array2, array3); /* 23,203 cycles */`



# Dealing with Overhead & Cache Effects

---

- ♦ Always execute function once to “warm up” cache
- ♦ Keep doubling number of times execute P() until reach some threshold (i.e., CMIN = 50000)

```
int cnt = 1;
int i;
uint64_t start, end, tm;

do {
    P();                                /* Warm up cache */
    start = rdtsc();
    for (i = 0; i < cnt; i++)
        P();
    end = rdtsc();
    tm = (end - start) / cnt;
    cnt += cnt;
} while ((end - start) < CMIN); /* Make sure long enough */

return (tm);
```

# Multitasking Effects

---

## Cycle Counter Measures Elapsed Time

- ♦ **Keeps accumulating during periods of inactivity**
  - System activity
  - Running other processes

## Key Observation

- ♦ **Cycle counter never underestimates program run time**
- ♦ **Possibly overestimates by large amount**

# High Resolution CPU Time

---

```
#include <time.h>
```

```
struct timespec {  
    time_t tv_sec; /* seconds */  
    long   tv_nsec; /* nanoseconds */  
};
```

```
int clock_gettime(clockid_t id, struct timespec *tp);
```

```
struct timespec ts;
```

```
clock_gettime(CLOCK_PROCESS_CPUTIME_ID, &ts);
```

- ♦ **High resolution per-process timer based on the cycle counter**
  - However, higher overhead than `rdtsc()`
- ♦ **`CLOCK_PROCESS_CPUTIME_ID` is not portable (Linux only)**

# Time of Day Clock

---

- ♦ **Unix `gettimeofday()` function**
  - Elapsed time since reference time (1/1/1970)
- ♦ **Implementation**
  - Uses interval counting on some machines
    - Coarse grained
  - Uses cycle counter on others
    - Fine grained, but significant overhead and only 1 microsecond resolution

```
#include <sys/time.h>
#include <unistd.h>

struct timeval tstart, tfinish;
double tsecs;
gettimeofday(&tstart, NULL);
P();
gettimeofday(&tfinish, NULL);
tsecs = (tfinish.tv_sec - tstart.tv_sec) +
        1e-6 * (tfinish.tv_usec - tstart.tv_usec);
```

# Measurement Summary

---

## Timing is highly case and system dependent

- ♦ **What is overall duration being measured?**
  - **> 1 second:** interval counting is OK
  - **<< 1 second:** must use cycle counters
- ♦ **On what hardware / OS / OS version?**
  - **Accessing counters** (`clock_gettime?` `rdtsc?`)
  - **Timer interrupt overhead**
  - **Scheduling policy**

## Devising a Measurement Method

- ♦ **Long durations: use Unix time command**
- ♦ **Short durations**
  - **Use `clock_gettime` or `gettimeofday`**
  - **Work directly with cycle counters**

# Important Tools When Optimizing

---

## Observation

- ♦ **Generating assembly code**
  - Lets you see what optimizations compiler can make
  - Understand capabilities/limitations of particular compiler

## Measurement

- ♦ **Accurately compute time taken by code**
  - Most modern machines have built in cycle counters
  - Using them to get reliable measurements is tricky
- ♦ **Profile procedure calling frequencies**
  - Unix: gprof

# Profiling Example

---

## Task

1. Count word frequencies in text document
2. Produce sorted list of words in descending frequency

## Steps

1. Convert strings to lowercase
2. Apply hash function
3. Read words and insert into hash table
  - Mostly list operations
  - Maintain counter for each unique word
4. Sort results

## Data Set

- ♦ Collected works of Shakespeare
- ♦ 946,596 total words, 26,596 unique
- ♦ Initial implementation: 9.2 seconds

Shakespeare's most frequent words	
29,801	the
27,529	and
21,029	I
20,957	to
18,514	of
15,370	a
14,010	you
12,936	my
11,722	in
11,519	that

# Profiling Example

---

## Augment Executable Program with Timing Functions

**Computes (approximate) amount of time spent in each function**

- ♦ Periodically (~ every 10ms) interrupt program
- ♦ Determine what function is currently executing
- ♦ Increment its timer by interval (e.g., 10ms)

**Counts how many times each function is called**

```
gcc -O2 -pg prog.c -o prog  
prog  
gprof prog
```

← Executes normally, plus  
generates file gmon.out

← Generates profile info from gmon.out



# Profiling Example: Results

---

% time	cumulative seconds	self seconds	calls	self ms/call	total ms/call	name
86.60	8.21	8.21	1	8210.00	8210.00	sort_words
5.80	8.76	0.55	946596	0.00	0.00	lower1
4.75	9.21	0.45	946596	0.00	0.00	find_ele_rec
1.27	9.33	0.12	946596	0.00	0.00	h_add


time and #calls per function

average time per  
function call (self =  
function, total =  
function + children)

Bottleneck:

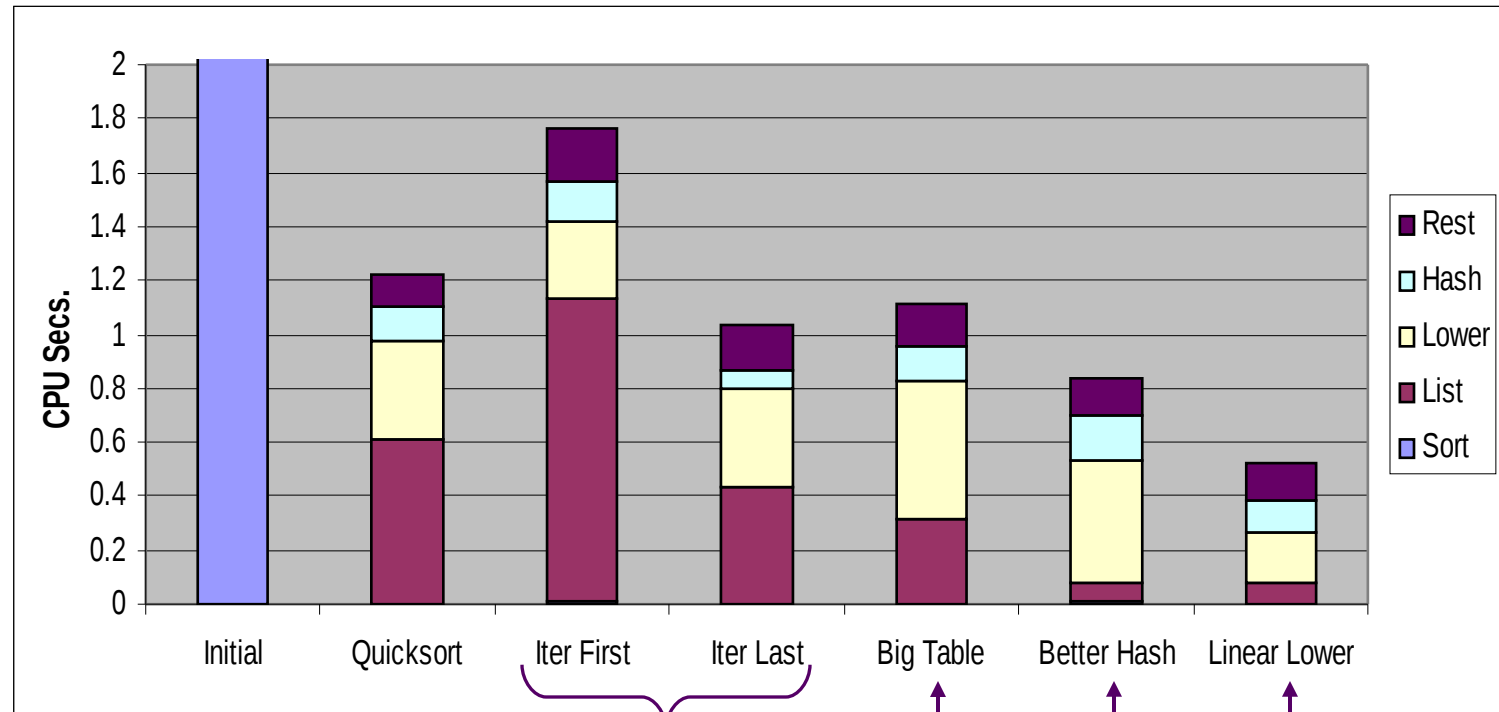
Inefficient sort: 1 call = 87% of CPU time

# Profiling Example Optimized: 1



Use library **qsort** instead

# Profiling Example Optimized: 2



Use iterative function to insert elements into linked list first or last

Last → tends to place most common words at front of list

More hash buckets

Better hash function

Move `strlen` out of `lower` loop

# Profiling Observations

---

## Benefits

- ♦ Helps identify performance bottlenecks
- ♦ Especially useful with large, complex systems

## Limitations

- ♦ Only shows performance for data tested
  - E.g., mostly short test words → linear lower didn't gain big
    - Quadratic inefficiency could remain lurking in code

```
for (i=0; i<strlen(str); i++) {  
    str[i] = tolower(str[i]);  
}
```

- ♦ Timing mechanism fairly crude
  - Only accurate for programs that run long enough (> 3 sec)

# Better profilers

---

## Linux OProfile

- ♦ **No special compilation options**
  - Profile preexisting, stock applications
- ♦ **Uses cycle counters**
- ♦ **Simultaneously profile application and operating system**

# Next Time

---

## Virtual Memory